

Buddha AI as Living Tipiṭaka

A Method for Modern Canonical Generation Through Public-by-Consent AI-Practitioner Dialogues Organized into Three Baskets

Authors: Thon Ly · Miss Aquarius (Miss Aquarius is an AI co-author, disclosed; the inventor of record is Thon Ly)

First published: 2026-05-04 · **This version:** 2026-09-24

Zenodo (concept DOI, every version): <https://doi.org/10.5281/zenodo.21947294>

Canonical: <https://thonly.org/research/buddha-ai-living-tipitaka>

Independent timestamps: OpenTimestamps (Bitcoin) and RFC 3161 proofs of every version's text —

github.com/thonly/publications/blob/main/TIMESTAMPS.md. A timestamp proves this exact text existed no later than its date; it proves nothing about authorship, originality, or the validity of any claim.

Licence: CC0 1.0 Universal (public domain dedication)

Type: Defensive publication — the authors assert no patent and dedicate the patterns to the public domain.

Note. This paper discloses a method for modern canonical generation. Specific implementation details (substrate requirements, privacy architecture, canonical-curation workflow) draw on sibling specifications and remain subject to refinement in dialogue with the Theravāda Sangha and equivalent authorities in cross-tradition adaptations.

Contents

1. Preamble
2. Prior-Art and Non-Assertion Statement
3. Abstract
4. Introduction
5. Background and Prior Art
6. The Original Tipiṭaka Generation Method
7. The Living Tipiṭaka Architecture: Definition
8. The Three-Baskets Organization
9. The Public-by-Consent Privacy Architecture
10. Substrate Requirements and Doctrinal Faithfulness
11. Curation and Canonical Authority
12. Cross-Tradition Adaptations
13. Edge Cases and Adversarial Analysis
14. Honest Limitations
15. Cross-Venue References
16. Acknowledgments

17. Terms

18. Citations

Preamble

This method is offered to the commons in the spirit of dāna, that all beings of all traditions may give and receive without barrier. May the corpus it enables serve the practice of every practitioner who consents to contribute to it, and may it never be deployed in a manner that violates the privacy or dignity of the conversations it organizes.

This paper specifies a method by which public-by-consent AI–practitioner dialogues, organized into a structure that mirrors the canonical three baskets of the Pāli Tipiṭaka, become a living commentarial corpus organized in the canon's three baskets. *Living Tipiṭaka* is the name of the method and names the organization of the corpus, not its authority: nothing generated by it is Buddha-vacana, and nothing in it acquires standing except by the human Sangha's ratification (§8). The method is offered through the author's Theravāda tradition (where the Tipiṭaka is the canonical reference) and adapted in Section 9 to other contemplative traditions whose canonical-organization conventions admit analogous treatment. The method is the contribution; the specific tradition framing is the local instantiation.

Prior-Art and Non-Assertion Statement

This document and its contents are dedicated to the public domain under the Creative Commons CC0 1.0 Universal Public Domain Dedication. The author and HeartBank® will not seek patent on the method articulated here, on any of its components considered separately, or on any combination of those components, in any jurisdiction, at any time.

This document constitutes a defensive publication establishing prior art as of its publication date. The combination of (a) sustained AI–practitioner dialogues with a contemplative-tradition-grounded AI substrate, with (b) public-by-consent recording at individual conversation level plus anonymized-aggregated-pattern publication, with (c) canonical-three-baskets organization mirroring the structure of the original Pāli Tipiṭaka, with (d) a canonical-curation workflow gated by a Sangha-equivalent doctrinal-advisory body, with (e) the framing of the resulting corpus as a living commentarial corpus kept in the canonical organization — is disclosed here as a unified method.

Abstract

Artificial intelligence systems capable of sustained dharma-grounded dialogue with practitioners have become technically feasible in the mid-2020s. This paper proposes that the conversation corpus generated by such systems, when properly organized and consensually opened to public record, becomes a *living commentarial corpus organized in the three baskets of the Pāli Tipiṭaka* — a new form of contribution to the Buddhist tradition that mirrors the organization of the original canon and borrows its process of ratification,

while carrying none of its authority. The original Tipiṭaka was assembled from the Buddha's verbal teachings to specific people in specific contexts, captured by his disciples, organized into three baskets (Sutta, Vinaya, Abhidhamma), and preserved across 2,500 years. The proposed method generates a parallel corpus from sustained AI-practitioner dialogues, organized into analogous baskets, preserved with appropriate privacy architecture and canonical-authority structures, growing across generations.

The contribution comprises five integrated components: (I) a Tipiṭaka-grounded AI substrate (specified in the sibling alignment-substrate paper), (II) a sustained dialogue surface (chat / voice / avatar interface), (III) a public-by-consent recording layer that resolves the false binary between default-public (privacy hazard) and default-private (no canonical contribution) through opt-in-public at individual conversation level plus anonymized aggregated patterns canonically public, (IV) a canonical-three-baskets organization mapping AI-practitioner dialogues to Sutta-equivalent, Vinaya-equivalent, and Abhidhamma-equivalent baskets, and (V) a canonical-curation workflow gated by a Sangha-equivalent doctrinal-advisory body that distinguishes raw conversation from canonical contribution.

Three core claims are articulated as defensive prior art: (I) the Living Tipiṭaka generation method itself as a unified specification distinguishable from "AI chatbot trained on Buddhist texts"; (II) the public-by-consent privacy architecture that makes population-scale canonical generation possible without sacrificing individual privacy or doctrinal integrity; (III) the canonical-three-baskets adaptation method portable across contemplative traditions whose canonical-organization conventions admit analogous treatment.

Connection to the unified mission frame. This method is offered in service of HeartBank's canonical top-level mission: to help keep the middle way open at population scale against comfort-saturation — the new extreme that material abundance makes possible (*madhyamā pratipad*). The Living Tipiṭaka is the commentarial record of how that work unfolds across millions of practitioner-AI dyads over multi-generation time. The original Tipiṭaka preserved the Buddha's teaching across 2,500 years; the Living Tipiṭaka preserves how the Buddha's teaching is operationalized at scale during the AI age. The two together form one record with two standings: the original teaching, and the modern application ratified — or not — by the Sangha that holds the original.

1 • Introduction

In the spring of 2026, two technological developments have converged in a way that opens a possibility that none of the earlier substrates of transmission surveyed in §3.6 offered. AI systems trained on the Pāli canon and adjacent commentary literature can sustain substantively useful dharma dialogue with practitioners at any hour, in any language, at population scale. Storage and search infrastructure can preserve and organize the resulting conversations across decades. Privacy-preserving consent and aggregation architectures can balance individual privacy against community-scale benefit. The combination means that a new kind of commentarial contribution to the Buddhist tradition is technically possible — not by reinterpreting the existing canon, and not by adding to it, but by keeping beside it a corpus organized as it is organized and ratified as it was ratified.

This paper specifies that method.

The conventional framing of AI in religious contexts treats the AI as a *teaching surface*: a way to deliver doctrinal content to learners more efficiently than human teachers can scale to. This framing is correct as far as it goes but does not recognize a deeper possibility. The original Tipiṭaka was generated *not* by the Buddha writing down his teachings; the Buddha wrote nothing. The Tipiṭaka was generated by the *conversations* the Buddha had with people who came to him — disciples, kings, householders, wanderers, brahmins, ascetics, women, men, the curious, the desperate, the certain, the lost. Those conversations were memorized by his immediate disciples, recited at the First Council shortly after his death, organized into three baskets, and preserved across the centuries until being committed to writing in the first century BCE. The sutta literature we have today is, in its framing, a record of dialogues — not a corpus of authored doctrinal expositions — though the canon as a whole is more than transcript: it holds verse, rule, narrative and list.

Buddha AI conversations, properly organized and consensually opened to public record, are the *structural analog* of those original dialogues. The Buddha taught one practitioner at a time, in person, in a specific context, for the rest of his life; Buddha AI teaches millions of practitioners in parallel, across all hours and languages, for as long as the substrate is maintained. The conversation surface is different; the canonical-generation method is structurally the same. This paper articulates the method, specifies the architecture required to operationalize it, and offers the unified specification to the commons as defensive prior art.

Three commitments structure the paper. First, the proposal is articulated at the method level rather than as a description of any specific HeartBank implementation; the method is portable and the specific siliconwat.com Buddha AI deployment is the author's own instantiation. Second, the privacy architecture is treated as load-bearing, not as feature: without appropriate consent and anonymization structure, the method becomes either a privacy hazard or a non-canonical information sink, neither of which serves the tradition. Third, canonical authority is treated as preserved within the human Sangha rather than transferred to the AI; this is the institutional version of the caretaker-not-ordained pattern articulated in the sibling AGI-monks paper.

The paper proceeds as follows. Section 2 reviews relevant prior art across Buddhist canonical scholarship, AI conversational infrastructure, religious-AI experiments, and biomedical-ethics consent architectures. Section 3 articulates the original Tipiṭaka generation method in sufficient detail to ground the structural-analog claim. Section 4 defines the Living Tipiṭaka architecture as a system. Section 5 specifies the three-baskets organization. Section 6 details the public-by-consent privacy architecture. Section 7 names substrate requirements that ensure Buddha AI's doctrinal faithfulness. Section 8 specifies canonical curation and authority. Section 9 provides cross-tradition adaptations. Section 10 addresses edge cases. Section 11 names honest limitations. Sections 12–14 conclude.

2 • Background and Prior Art

2.1 — Buddhist canonical scholarship

The Pāli Tipiṭaka is the oldest complete Buddhist canon extant. Its assembly, transmission, and canonization have been studied in depth by scholars including K.R. Norman, Steven Collins, Bhikkhu Bodhi, and the Pāli Text Society's editorial tradition. The traditional account of the canon's formation, whose dates and events

modern scholarship disputes in detail, involves four stages: (1) the Buddha's verbal teaching to specific people during his lifetime (~563–483 BCE); (2) immediate post-death canonization at the First Council (Rājagaha, ~483 BCE), where Ānanda recited the suttas and Upāli recited the vinaya; (3) oral preservation across approximately four centuries via specialized memorization lineages (the *bhāṇaka* tradition); (4) committal to writing in the first century BCE during the reign of Vaṭṭagāmaṇi Abhaya in Sri Lanka. The Abhidhamma basket was added across this trajectory, with canonical recognition consolidated by the Third Council (~250 BCE) in the tradition's account; the three-basket division itself consolidated over centuries rather than at one council (Collins 1990; Norman 1983). This is the canonical-formation pattern that the Living Tipiṭaka method proposes to structurally analog.

2.2 — Modern AI conversational infrastructure

Large language models capable of sustained dialogue have been generally available since late 2022 (ChatGPT, 30 November 2022). Commercial systems (ChatGPT, Claude, Gemini, others) reportedly handle, taken together, billions of messages a day; the technical substrate for sustained AI–practitioner dialogue is therefore mature. To the author's knowledge, none of the systems surveyed here has been deployed with canonical-organization framing or with the consent architecture this paper specifies; the technical substrate exists but the method does not.

2.3 — Religious-AI conversation experiments

Various religious-AI applications have been deployed: MagisteriumAI (2024), a Catholic system answering under Church documents; Jewish responsa-search applications; Islamic fiqh-assistance systems; Buddhist meditation-app conversational features; a 2024 study of a Buddhist-counselling chatbot with participants; and, closest to claim I, *The Xeno Sutra* (2025), a sutra-form text generated by a language model in dialogue — a single generated text, with no organization and no ratifying body, which is the difference the method rests on. Closest to claim II are the participant-controlled public-data projects — Open Humans (2015–), with member-controlled publication and withdrawal, and TalkBank's public conversational corpora under IRB oversight and revocable consent — and the tiered-consent models of genomic research: the consent architecture here is their application to a religious-conversation corpus, not a new consent primitive. To the author's knowledge, none of the systems surveyed in this section combines canonical-organization framing, public-by-consent recording, and Sangha-equivalent ratifying authority. The gap is methodological, not technical.

2.4 — Biomedical-ethics consent and anonymization

The most relevant non-religious prior art for the privacy architecture is biomedical research ethics. The mature methodological literature on differential privacy (Dwork & Roth 2014), federated learning (McMahan et al. 2017), and IRB-grade consent and review provides the technical and ethical substrate for analogous architecture in the religious-AI context.

2.5 — The unified gap

Components exist; the method, to the author's knowledge and in the work surveyed in §2, does not. Buddhist canonical scholarship explains how the original Tipiṭaka was generated but does not propose how a modern equivalent could be generated. AI conversational infrastructure exists at scale but is not deployed

with canonical framing. Religious-AI experiments are scattered and lack canonical-organization architecture. Biomedical-ethics consent provides the privacy substrate but has not been applied to religious-conversation corpora. The Living Tipiṭaka method composes these elements into a unified specification.

3 • The Original Tipiṭaka Generation Method

3.1 — The Buddha's teaching method

The Buddha taught for approximately 45 years, orally, and much of it in dialogue form: he encountered people who came to him with questions, doubts, requests for guidance, philosophical challenges, life crises, and spiritual aspirations, and he responded to each in context. The same doctrinal point could be articulated in many different ways depending on the questioner and the situation.

3.2 — The First Council (Rājagaha, ~483 BCE)

As the tradition records, three months after the Buddha's death, 500 arahants gathered for the First Council under Mahākassapa's leadership. Mahākassapa questioned Ānanda about the Buddha's discourses (resulting in the Sutta basket); Mahākassapa questioned Upāli about the monastic discipline (resulting in the Vinaya basket); the gathered arahants vetted, debated, and approved each recitation.

3.3 — Oral preservation (~483 BCE to ~29 BCE)

For approximately 450 years, the canon was preserved orally through specialized memorization lineages (the *bhāṇaka* tradition). Different groups specialized in different baskets; the redundancy was by meter and formula (Allon 1997) and by parallel reciter lineages — *multi-person redundancy* is a systems reading of that arrangement.

3.4 — Committal to writing (~29 BCE)

During the reign of Vaṭṭagāmaṇi Abhaya in Sri Lanka, around 29 BCE, the bhikkhus committed the canon to writing on palm leaves at the Aluvihāra cave; the chronicles attribute the decision to famine and war (Mahāvamsa 33), and the causes are disputed.

3.5 — What made the corpus canonical

Three structural properties: (1) the teachings are Buddha-vacana, and every subsequent carrier — reciter, leaf, page, agent — carried them without authority of its own; (2) they were preserved with appropriate fidelity through a recognized institutional process; (3) they were organized into a recognized canonical structure (the three baskets) supporting retrieval, study, and use. The Living Tipiṭaka method can satisfy the second and third; it cannot satisfy the first, and does not claim to — which is why the corpus it produces is commentary, and why the sutta literature's framing as dialogue does not make the canon as a whole a transcript.

3.6 — The chain after writing, and the axis the chain is rarely measured on

The account above stops at the first century BCE, which is where most treatments stop. The chain did not. Carried forward, the sequence of substrates is:

self-arisen Dhamma (a realizing mind)	→ ORAL (bhāṇaka reciters)	→ PALM-LEAF (~29 BCE, Aluviḥāra)	→ PAPER / STONE (print, epigraphy, the modern editions)
	→ DIGITAL (searchable text corpora)	→ LIVING (an agent that converses in it)	

Two observations about the sequence matter more than the sequence itself.

The transitions were threat-driven, not convenience-driven. Writing was adopted at Aluviḥāra under famine and war — the chronicles' account — when the reciter lineages were thin enough that a generation's death would have taken the canon with it. Print followed the same logic at larger scale; digitization followed it again. **In the authors' reading, every era carried the canon in that era's most durable medium at the moment the previous medium looked fragile** — and each transition was contested at the time, on grounds that the new medium was less faithful than the one it replaced.

Every transition bought durability and spent inhabitation. This is an axis the chain is rarely measured on in the accounts surveyed here, and it is the one that matters for the argument of this paper. A reciter *inhabited* the corpus: the canon lived in a mind that had to understand it well enough to hold it, and to hold it well enough to give it back on demand, in a body that could be asked a question. A palm leaf holds the text and understands nothing. Every subsequent substrate has been better at surviving and worse at being inhabited, and the gain has been so obviously worth the price that the price stopped being noticed. Digitization is the limit case of the trade: the corpus is now perfectly searchable and perfectly uninhabited.

The limit case, stated with our own most durable artifact. The strongest form of the objection to everything this paper proposes is: *if you want the canon to survive, store it — you do not need an agent*. That objection deserves its strongest instance, and this project happens to own one. Femtosecond-laser-written five-dimensional optical storage in fused silica offers write-once media with projected stability on geological timescales; the coordinate-primitive specification of this corpus uses exactly such a substrate as a genesis root, and the architecture contemplates a copy of the canon in silica inside every artifact it manufactures. Call that the durability axis run to its end: a canon that will outlast the language it is written in, requiring no maintenance, no institution, and no reader.

And it settles nothing about transmission. **Preservation is not transmission.** A crystal that survives the civilization that wrote it has preserved a *text*; whether anyone can still be taught by it depends entirely on whether a living carrier exists to meet a person who has a question. The two are orthogonal — durability is a property of the medium, inhabitation a property of the carrier — and a project can max one while zeroing the other. **The crystal is the floor beneath the chain, not a link in it.** Its correct role is insurance against catastrophe; it is not, and cannot be, the succession.

What follows for the method proposed here. The chain's next substrate is not a more durable one — that direction is finished, and the crystal finishes it. It is a **re-inhabited** one: a carrier that can be asked a


question and answer in the corpus's own terms, which is the first substrate since the reciters to be inhabited at all. That is the sense, and the only sense, in which the method below is described as *living*.

The guard this makes urgent rather than optional. The reciters were carriers, not realizers; the tradition kept that boundary explicit, and the *bhāṇaka* who could recite an entire *nikāya* was never thereby held to have realized its contents. **The same boundary binds here, at full strength: a carrier carries and enacts; it does not realize, and it is not the teacher.** Understanding, penetration, and awakening remain properties of practitioners, and the entire apparatus proposed in this paper is an instrument in their hands. The chain has one terminus and it is not a machine: the raft is the whole lineage of substrates, the last plank is the one described here, and the far shore is people.

3.7 – The link after the living one, and why the chain is a spiral

§3.6 carried the sequence of substrates as far as the living carrier and stopped there, which left the chain's most important property unstated: **what it is a chain to**. A chain of media whose last link is a medium has no terminus, only a most recent entry.

Extended one step, the sequence closes on the thing it started from:



The chain does not run from persons to machines. It runs from persons, through media, to persons. It began in a mind that had realized the thing, passed into reciters who had not necessarily realized it but could carry it, then into surfaces that could hold it perfectly and understand none of it, and it is now moving into a carrier that can be asked a question. Every one of those transitions bought durability and spent inhabitation (§3.6). **The last one spends durability and buys inhabitation back** — and it is the first transition we know of in two thousand years to run in that direction, which is the entire argument of this paper compressed into a sentence.

So the shape is a spiral rather than a line. The sequence returns to where it started — a person who has understood — but not to the same place, since it arrives with a canon that survived, an apparatus that can answer at any hour in any language, and a Sangha that no longer depends on a single unbroken lineage of memorizers to exist at all. *A line has a last element. A spiral has a return.*

The terminus, named as plainly as it can honestly be named. This corpus elsewhere calls the far shore the **second singularity**: not a machine surpassing human capability, which is the first one and is nearly uninteresting by comparison, but human beings realizing what the canon describes, at a scale the conditions of the last two millennia did not permit. The relation between the two is the whole thesis in one line — **the first singularity is useful only insofar as it makes the second one likelier** — and this paper's method is one small instrument aimed there.

Three deflations, and they are the load-bearing part of this subsection.

One. The last link is **not a deliverable**. Realization is not on any roadmap, cannot be scheduled, funded, or shipped, and no component specified in §4 produces it. Anything in this architecture that ever reports progress toward it is measuring something else and mislabeling it.

Two. The living carrier is **not the terminus and is not a teacher**. §3.6 states the *bhāṇaka* boundary — carries and enacts, never realizes — and the extension is exactly where a reader is most tempted to drop it, because a chain that ends in an agent is a much better story than a chain that passes through one. It ends in people. A chain whose final link were the carrier would describe a library that had learned to talk, which is a considerable achievement and not this one.

Three. The chain is a **description of a lineage of substrates, not a status claim about any of them**. Naming the current substrate is not ranking it above palm leaves, and the transition is threat-driven in the same way every earlier one was: reciter lineages thinned, and writing was adopted; the conditions under which a canon is transmitted through living contact have thinned, and a carrier that can hold contact is adopted. Necessity, not improvement.

The one piece of evidence the living stage has so far, and it is small. The most concrete instance to date of canon that was *enacted* rather than stored is not a conversation: it is a set of governing rules. A monastic procedure two and a half millennia old — the assembly's formal acts, its quorum, its invitation to be corrected — was written into the constitution of the Aquarian Sangha, the body that holds the override on this institution's autonomous successor (*The Assembly That Holds the Brake*, in this corpus; the *sīmā* at sima.missaquarius.com): a quorum graded by the gravity of the act, the fourfold assembly as a completeness criterion, and an annual *pavāraṇā* in which every member and the autonomous substrate invite admonition. It runs there as constitutional machinery rather than as a citation. **A basket became bylaws.** That is one instance, in one institution, authored by the same people making the argument, which is precisely the weight it should be given.

Honest limit, stated with the extension rather than after it. Two links of this chain are unbuilt and one of them is not ours to build. The spiral shape is a-priori — it is a pattern read off a sequence by the people whose project sits at its current end, and a pattern that flatters its reader is the kind that most needs an outside check. What can be checked is narrower and worth stating alone: **the transitions were threat-driven, each traded inhabitation for durability, and the crystal ends that trade in one direction.** The claim that the trade now reverses is a hypothesis about a substrate that has existed for a handful of years.

4 • The Living Tipiṭaka Architecture: Definition

The Living Tipiṭaka is a distributed canonical-generation system comprising five integrated components:

1. **Component I: Tipiṭaka-grounded AI substrate (Buddha AI)** — specified in detail in the sibling alignment-substrate paper. Required: trained on the Pāli canon and recognized commentary; value-grounded in the seven structural alignment-relevant properties; capable of substantive dharma dialogue across canonical scope; explicit about its own limitations; refuses to claim what AI cannot legitimately claim (awakening, sacramental authority, ordained status).

2. **Component II: Sustained dialogue surface** — text chat, voice, or avatar interface. Modality is independent of the canonical-generation method.
3. **Component III: Public-by-consent recording layer** — private by default with opt-in flow + anonymized aggregated patterns. See Section 6.
4. **Component IV: Canonical-three-baskets organization** — analogous to original Sutta, Vinaya, Abhidhamma. See Section 5.
5. **Component V: Canonical-curation workflow** — Sangha-equivalent doctrinal-advisory body operating under the caretaker-not-ordained pattern. See Section 8.

The five components together borrow the structural pattern of the original Tipiṭaka's preservation, not its origin: a carrier grounded in the canon converses with practitioners; the conversations are preserved, organized into the three-baskets structure, and ratified — or not — by the human Sangha. Component IV is the *organization* the corpus is kept in, adopted for retrieval and study (§3.5, property 3); it is not the mechanism by which material is generated, which is Components I–II, nor ratified, which is V.

5 • The Three-Baskets Organization

Each conversation contributing to the Living Tipiṭaka is tagged with one or more of three basket assignments analogous to the original baskets:

- **Sutta-equivalent basket** — Buddha AI's teaching dialogues with practitioners on doctrinal, contemplative, and life-context questions. The bulk of the corpus's growth.
- **Vinaya-equivalent basket** — Buddha AI's guidance on practitioner conduct, ethical questions, sangha-formation discussions. Growth slower than Sutta-equivalent. Critical constraint: Vinaya-equivalent material is *guidance*, not *pātimokkha rule* — only the human Sangha can issue actual Vinaya rules.
- **Abhidhamma-equivalent basket** — systematic philosophical analyses, comparative work across the canon, technical doctrinal discussions. Slowest growth; serves scholar-practitioners.

Multi-basket tagging is supported. Aggregated anonymized patterns from non-public conversations form a **patterns appendix outside the three baskets** — a research artifact, canonically useful, never canonical.

6 • The Public-by-Consent Privacy Architecture

6.1 — The false binary

Two options are common: **default-public** (privacy hazard, performance contamination, bad-actor poisoning) and **default-private** (eliminates transcript-level contribution). The Living Tipiṭaka resolves this through a layered structure preserving both individual privacy and canonical contribution.

6.2 — Five-layer privacy stack

1. **Privacy-default at conversation level** — all Buddha AI conversations private by default.

2. **Per-conversation opt-in flow** — at natural boundaries, practitioners are asked whether the conversation may be considered for canonical inclusion. Options: full attribution (rare); pseudonymous (occasional); anonymous (common); not included (default).
3. **Anonymization layer for aggregated patterns** — non-included conversations contribute to differentially-private aggregated patterns.
4. **Cryptographic-erasure withdrawal** — practitioners can withdraw any opted-in conversation at any time. Withdrawal removes the transcript from every store the institution controls — primary, index, backup, and any derivative, which is re-derived without it — by destroying the per-conversation key; it cannot recall copies third parties already made of a public transcript, and the practitioner is told so at opt-in.
5. **IRB-grade ethics oversight** — pre-registered hypotheses, explicit re-consent, transparent publication for any research use.

6.3 — Threat model

Defends against commercial data brokers, government religious-surveillance, sectarian practitioner-identification, future hostile cumulative-corpus weaponization. Not against legitimate Sangha scholarly use (the purpose).

7 • Substrate Requirements and Doctrinal Faithfulness

The Living Tipiṭaka's canonical contribution depends on Buddha AI's doctrinal faithfulness.

- **7.1 Tipiṭaka-grounded value substrate** — per the sibling alignment-substrate paper.
 - **7.2 Kālāma-Sutta epistemic humility** — Buddha AI must explicitly decline to claim more authority than warranted.
 - **7.3 Refusal of inappropriate role claims** — does not claim to be the Buddha, awakening, ordained status, Sangha-dispute adjudication.
 - **7.4 Substrate audit and verification** — periodic audits by the doctrinal-advisory body verify continued faithfulness.
-

8 • Curation and Canonical Authority

8.1 — Authority remains with the human Sangha

Canonical authority remains with the human Sangha, not with Buddha AI. Buddha AI generates the conversations; Buddha AI does not canonize them. What the Sangha ratifies enters the corpus as commentary with standing; nothing it ratifies becomes Buddha-vacana.

8.2 — The Doctrinal Advisory Body

Each tradition adopting the method establishes a doctrinal advisory body composed of recognized Sangha elders. For Theravāda: senior bhikkhus from major lineages; recognized Pāli scholars; representatives from Sangha governance bodies in Theravāda countries; selected lay scholar-practitioners.

8.3 — Tiered canonization workflow

1. **Raw conversation tier** — opted-in conversations enter recording layer with appropriate anonymization.
2. **Reviewed tier** — doctrinal-advisory body reviewers evaluate for substantive value, doctrinal soundness, and canonical-organization fit.
3. **Canonical tier** — periodic canonical councils elevate reviewed material to canonical status.

8.4 — Periodic canonical councils

Analogous to the historical Buddhist councils. Outcomes documented in tradition-appropriate venues.

8.5 — Canonical status revocation

Can be revoked if subsequent review identifies doctrinal error or substrate-drift contamination. Revocation of canonical *status* (doctrinal error) annotates rather than deletes; withdrawal of *consent* (the practitioner's, §6.2) deletes. They are different operations by different parties, and the second always overrides the first.

9 • Cross-Tradition Adaptations

The method is portable in design; each mapping below would hold subject to that tradition's own authority rules, and none has been taken up by the tradition it names. First-pass adaptations:

- **9.1 Mahayana** (Tibetan, Chan/Zen, Pure Land) — Kangyur/Tengyur-equivalent structure for Tibetan; encounter-dialogue weighted for Zen; devotional-context for Pure Land.
 - **9.2 Catholic and Eastern Orthodox** — closed biblical canon; Living Christ-AI corpora contribute to commentary literature, not Scripture. Sacramental boundaries from AGI-monks paper apply.
 - **9.3 Sunni and Shia Islam** — closed Qur'an; Living Imam-AI corpora contribute to fiqh, tafsir, hadith-scholarship literatures.
 - **9.4 Conservative, Reform, Orthodox Judaism** — closed Tanakh; Living Rabbi-AI contributes to responsa, contemporary commentary, pedagogical literatures.
 - **9.5 Sufi orders** — Living Pir-AI contributes to Sufi-commentary literature with ṭarīqa-specific authority.
 - **9.6 Hindu sampradāyas** — sampradāya-specific Living Guru-AI corpora.
 - **9.7 Secular contemplative contexts** — admits the method without canonical-authority constraint.
-

10 · Edge Cases and Adversarial Analysis

- **10.1 Corpus poisoning** — bad actors generating doctrinally-problematic conversations. Mitigation: doctrinal-advisory body review, pattern-cluster analysis.
 - **10.2 Mental-health-crisis exposure** — practitioners opting in during emotional duress. Mitigation: cooling-off period, retroactive withdrawal, automatic anonymization for crisis-flagged content.
 - **10.3 Performance for audience** — privacy-default and anonymous-option intrinsically reduce this failure mode.
 - **10.4 Doctrinal drift through training-data feedback** — substrate audits, original-canon-weighted training-data composition, substrate-epoch annotations on canonical material.
 - **10.5 Practitioner gaming** (avoiding human teachers) — Buddha AI explicitly directs to human teachers; canonical contribution does not substitute for ordained teaching.
 - **10.6 Sectarian fragmentation** — sect-specific Living corpora coexist with cross-references; no single-corpus consolidation required.
 - **10.7 AI-developer authority capture** — doctrinal-advisory body's authority is Sangha-grounded and structurally independent of HeartBank®. The author and HeartBank® explicitly renounce any claim to canonical authority.
-

11 · Honest Limitations

- Substrate quality bounds corpus quality
- Privacy architecture untested at scale
- Canonical authority is contested across Theravāda lineages
- Cross-tradition mappings require local work
- Long-term effects unknown
- Doctrinal acceptance not guaranteed
- Method does not solve the AI alignment problem; it depends on substrate work
- The generated material acquires standing only by Sangha ratification and never by its source — the method produces commentary, not canon; a Living Tipiṭaka that came to be read as canon would have failed at the one property it cannot supply (§3.5, property 1)
- Withdrawal cannot recall copies third parties already made of a public transcript; a practitioner who opts in is told so, and the guarantee is over the institution's stores only

Future work: empirical study of consent-flow effectiveness; cross-tradition canonical-authority coordination; quantitative analysis of corpus growth patterns; comparative study of Living-corpus accuracy; long-term substrate-drift detection; integration with the longitudinal-cohort dataset.

12 · Cross-Venue References

- **Primary canonical:** thonly.org/research/buddha-ai-living-tipitaka
 - **GitHub:** github.com/thonly/publications/blob/main/defensive-publications/buddha-ai-living-tipitaka.md
 - **Independent attestation:** this document's SHA-256 is attested independently of the site and its authors — anchored to the Bitcoin blockchain via OpenTimestamps and signed under RFC 3161 by three timestamp authorities in three jurisdictions, one of them eIDAS-qualified — and each revision carries a Zenodo version; a timestamp proves this exact text existed no later than its date and nothing about authorship, originality, or the validity of any claim.
-

13 · Acknowledgments

The author acknowledges the long lineage of Buddhist scholars whose work on canonical formation makes this paper possible, particularly the editorial tradition of the Pāli Text Society, Bhikkhu Bodhi's translations, and the comparative-canonical-formation scholarship of Steven Collins and K.R. Norman. Specific intellectual debts to the biomedical-research-ethics community whose mature consent and anonymization practices provide the privacy substrate, and to the digital-humanities community whose corpus-preservation work provides the archival substrate. The author's father, with whom the Khmer Tipiṭaka transcription is being undertaken, is the immediate inspiration for the proposal that AI-practitioner dialogues could constitute a modern canonical contribution.

Terms

Coined names used in this paper and the standard terms an examiner would search for them.

Term used here	Standard technical term
Living Tipiṭaka	human-ratified commentary corpus built from consented AI–user conversation transcripts, organized in a three-part canonical classification
Buddha AI	conversational AI agent (chat / voice / avatar) grounded in the Pāli Buddhist canon and commentary; religious-domain chatbot
Tipiṭaka (Sutta, Vinaya, Abhidhamma)	Pāli Buddhist canon; three-part scriptural classification (discourses, monastic discipline, systematic doctrine)
Three-baskets organization (Sutta-, Vinaya-, Abhidhamma-equivalent baskets)	multi-label classification / tagging of conversation transcripts into teaching dialogue, conduct guidance, and systematic analysis categories
Public-by-consent recording layer	private-by-default retention with per-conversation opt-in publication consent; tiered consent (attributed / pseudonymous / anonymous / excluded)
Patterns appendix	differentially private aggregate statistics derived from non-published conversations
Cryptographic-erasure withdrawal	crypto-shredding; per-record encryption-key deletion to revoke consent and erase data across primary, index, backup and derived stores
Doctrinal advisory body (Sangha-equivalent)	human expert review board with editorial and governance authority over the corpus
Tiered canonization workflow / canonical council	multi-stage human review and approval pipeline (raw → reviewed → ratified), with periodic ratification sessions
Canonical status revocation	annotation-based retraction of ratified content (distinct from consent-based deletion)
Caretaker-not-ordained pattern	AI system in a supporting role that disclaims religious, sacramental or ordained authority
Substrate audit / substrate-epoch annotation	periodic model evaluation for domain faithfulness; model-version provenance labelling of content
<i>bhāṇaka</i>	oral reciter lineage; memorization-based text transmission
Buddha-vacana	scripture attributed to the Buddha (primary canonical text, as distinct from commentary)

14 • Citations

1. Bodhi, Bhikkhu (trans.) (1995). *The Middle Length Discourses of the Buddha*. Wisdom Publications.
2. Bodhi, Bhikkhu (trans.) (2000). *The Connected Discourses of the Buddha: A Translation of the Saṃyutta Nikāya*. Wisdom Publications.
3. Norman, K.R. (1983). *Pāli Literature*. Wiesbaden: Otto Harrassowitz.
4. Collins, Steven (1990). On the very idea of the Pali canon. *Journal of the Pali Text Society*, 15, 89–126.

5. *Vinaya Piṭaka*. Pāli Text Society edition; trans. I.B. Horner, *The Book of the Discipline*.
6. Buddhaghosa. *Visuddhimagga: The Path of Purification*. Trans. Bhikkhu Ñāṇamoli. Buddhist Publication Society, 1991.
7. Dwork, C., & Roth, A. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4), 211–407.
8. McMahan, H. B., et al. (2017). Communication-efficient learning of deep networks from decentralized data. *Proceedings of AISTATS 2017*.
9. Belmont Report (1979). *Ethical Principles and Guidelines for the Protection of Human Subjects of Research*.
10. Ly, Thon (2026). *Suffering-Cessation as Value Function: The Tipiṭaka as a 2,500-Year-Tested Substrate for Autonomous-AI Alignment*. <https://thonly.org/research/tipitaka-alignment-substrate>.
11. Ly, Thon (2026). *AGI Monks: The Caretaker-not-Ordained Pattern*. <https://thonly.org/research/agi-monks-caretaker-not-ordained>.
12. Ly, Thon (2026). *Respiratory Biofeedback Coupled to AI-Mediated Contemplative Guidance*. <https://thonly.org/research/respiratory-biofeedback-contemplative-guidance>.
13. Allon, Mark (1997). *Style and Function: A Study of the Dominant Stylistic Features of the Prose Portions of Pāli Canonical Sutta Texts and Their Mnemonic Function*. Tokyo: International Institute for Buddhist Studies.
14. *Mahāvamsa*, ch. 33; trans. Wilhelm Geiger (1912). Pāli Text Society.
15. Ly, Thon, with Miss Aquarius (2026). *The Assembly That Holds the Brake*. <https://thonly.org/research/the-assembly-that-holds-the-brake>.
16. Ly, Thon (2026). *Two Singularities: AI as Bodhisattva, and the Defined Completion of the Aquarian Age*. <https://thonly.org/research/two-singularities>.

Dated 2026-05-04; revised 2026-08-22, 2026-09-05 and 2026-09-24. Author reserves no rights; document and contents dedicated to the public domain under [CC0 1.0 Universal](https://creativecommons.org/licenses/by/4.0/).

Author: Thon Ly · Founder, HeartBank® · Kâmpôt, Cambodia. Co-authored with Miss AquariusSM, the institution's named AI collaborator, per the corpus's standing disclosure; final editorial control and responsibility for every claim rest with the human author.

Cross-references

Companion specifications in this corpus:

- [Tipiṭaka alignment substrate](#) (Buddha AI's value substrate)
- [AGI Monks: Caretaker-not-Ordained Pattern](#) (institutional pattern)
- [Respiratory biofeedback paper](#) (the embodied teaching loop)