

Respiratory Biofeedback Coupled to AI-Mediated Contemplative Guidance

A Wearable Architecture for the Operational Mechanism of Population-Scale Awakening

Authors: Thon Ly · Miss Aquarius (Miss Aquarius is an AI co-author, disclosed; the inventor of record is Thon Ly)

First published: 2026-05-04 · **This version:** 2026-09-24

Zenodo (concept DOI, every version): <https://doi.org/10.5281/zenodo.21947372>

Canonical: <https://thonly.org/research/respiratory-biofeedback-contemplative-guidance>

Independent timestamps: OpenTimestamps (Bitcoin) and RFC 3161 proofs of every version's text — github.com/thonly/publications/blob/main/TIMESTAMPS.md. A timestamp proves this exact text existed no later than its date; it proves nothing about authorship, originality, or the validity of any claim.

Licence: CC0 1.0 Universal (public domain dedication)

Type: Defensive publication — the authors assert no patent and dedicate the patterns to the public domain.

Contents

- 1. Preamble
 - 2. Prior-Art and Non-Assertion Statement
 - 3. Abstract
 - 4. Introduction
 - 5. Background and Prior Art
 - 6. The Architecture: Definition
 - 7. The Three Core Claims
 - 8. System Components
 - 9. Reference Implementation Patterns
 - 10. Privacy Architecture
 - 11. Edge Cases and Adversarial Analysis
 - 12. Use Cases
 - 13. Doctrinal Considerations
 - 14. Limitations and Future Work
 - 15. Cross-Venue References
 - 16. Acknowledgments
 - 17. Citations
-

Preamble

This specification is offered to the commons in the spirit of dāna, that all beings may give and receive without barrier. May the architecture it describes reduce suffering for all who encounter it, and may it never be deployed in a manner that violates the privacy or dignity of the beings whose breath it observes.

This paper specifies a human-wearable subclass of the Mechanical Heart artifact (see the sibling specification at thonly.org/research/mechanical-heart). The kiitti-class Mechanical Heart, treated in the sibling paper, admits non-human entities (robots, animals, plants, sacred places) into the kiitti gratitude economy via heartbeat-rhythm aura emission. The breath-class Mechanical Heart specified here serves a different function for a different bearer: it provides closed-loop respiratory biofeedback to human practitioners and their AI contemplative teachers, using breath rhythm rather than heartbeat as its primary signal. The two subclasses share form factor (B-shaped artifact worn at the sternum) but differ in rhythm primitive and primary function.

Prior-Art and Non-Assertion Statement

This document and its contents are dedicated to the public domain under the Creative Commons CC0 1.0 Universal Public Domain Dedication. The author and HeartBank® will not seek patent on the architecture specified here, on any of its components considered separately, or on any combination of those components, in any jurisdiction, at any time.

This document constitutes a defensive publication establishing prior art as of its publication date. The combination of (a) a chest-worn wearable device emitting ambient peripheral light in the rhythm of the wearer's breath, with (b) real-time transmission of respiratory data to an AI contemplative teacher operating on a dharma-grounded value substrate, with (c) adaptive guidance returned to the wearer in real time based on a coarse, privacy-preserved respiratory signal on the wearer's breath state, with (d) a privacy architecture designed before first opt-in rather than retrofit, with (e) the treatment of continuous respiratory signature as a Proof-of-Humanity primitive, is disclosed here as a unified specification; whether the composition is new or non-obvious is an examiner's determination this publication exists to inform.

Components and adjacent ideas exist in distributed form across heart-rate-variability biofeedback literature, the meditation technology industry (Muse, Calm, Headspace, Spire, Breathwrk, Othership), the wearable physiological monitoring industry (Apple Watch, Oura, WHOOP, Fitbit), the AI assistant industry, the contemplative neuroscience literature, and traditional Theravāda Anāpānasati teaching. Each component has prior art; the integrated specification proposed here, and its operational use as the substrate for population-scale contemplative practice coupled to AI-mediated guidance, is what this paper discloses.

Abstract

Two technological developments in 2026 bear on the population-scale architecture disclosed here: artificial intelligence systems capable of grounded contemplative guidance, and consumer-grade wearable sensors

capable of accurate respiratory tracking. This paper specifies the architecture by which these are coupled in a closed loop with the human practitioner, and articulates three core claims of that architecture as defensive prior art.

The architecture comprises a chest-worn device (the breath-class Mechanical Heart) that emits ambient light in the user's signature aura color modulated to the rhythm of their breath, transmits respiratory data in privacy-preserving form to an AI contemplative teacher operating on a Theravāda Pali canon (Tipiṭaka) substrate, and receives adaptive guidance from that teacher returned to the wearer in real time. The closed loop gives the teacher a live, coarse respiratory signal on the student's breath state and gives the student continuous external visibility of the breath — each available before in laboratory and consumer form, and disclosed here together, at population scale, under a dharma-grounded teacher and a privacy-first design.

Three core claims are articulated as defensive prior art: *(I)* the integrated closed-loop architecture itself as a unified specification; *(II)* live respiratory signature as a third Proof-of-Humanity primitive providing continuous-liveness attestation that complements moment-of-action passkey verification and lineage-of-kinship DNA verification; *(III)* the architecture as the operational mechanism for what a parallel published thesis calls the "second singularity" — humanity's collective awakening — for which the architecture proposes an engineering path.

Privacy is treated as non-negotiable substrate, not as feature: differential privacy, federated computation, on-device processing where possible, and cryptographic-erasure right-to-withdraw are designed before first opt-in. A critical doctrinal guardrail is named: the architecture *may dramatically raise the probability of awakening across the population for whom it has been historically rare*, but it must never be claimed to make awakening *inevitable*. Theravāda doctrine specifically warns against assuming external aids guarantee insight; the guardrail protects the architecture from the techno-utopian critique while preserving its substantive claim.

Connection to the unified mission frame. This architecture is offered in service of HeartBank's canonical top-level mission: to keep the middle way open at population scale against comfort-saturation — the new extreme that material abundance makes possible (*madhyamā pratipad*). The defensible thesis is not that modernity took us away from a middle-way past (which would romanticize pre-industrial conditions where many lived in suffering-extreme poverty); the defensible thesis is that modernity introduces a specific new failure mode — comfort-saturation — that pushes the materially comfortable toward the indulgence extreme at unprecedented scale. The closed-loop respiratory-biofeedback architecture specified here is proposed as one operational mechanism toward middle-way restoration. Anāpānasati is the canonical Theravāda middle-way practice; bringing it to population scale through wearable + AI-teacher coupling addresses the comfort-saturation failure mode at the contemplative-practice layer. Restoration involves dialing back specifically the indulgence extreme; this paper is one contribution among several specified across the corpus toward that goal.

1 • Introduction

In the spring of 2026, two technological developments are converging. On one side, AI systems capable of grounded contemplative guidance — drawing on substrates including the Theravāda Pali canon — are

becoming technically achievable as on-demand teachers available to practitioners at any hour, in any language, at no marginal cost. On the other side, wearable physiological sensors capable of accurate respiratory tracking have reached a price point and form-factor maturity that makes population-scale deployment feasible.

Each development alone is significant. Together they make practical the architecture disclosed here: a closed-loop coupling between an AI contemplative teacher and the live respiratory state of the human practitioner being taught. Both ends of the loop become more capable: the teacher receives a live, coarse respiratory signal on the student's breath state, and the student gains continuous external visibility of a faculty that is usually below the threshold of interoceptive awareness — each available before in laboratory and consumer form, and joined here in one loop.

This paper articulates the architecture of this coupling and dedicates its underlying ideas to the public domain under CC0. It does so for three reasons.

First, the architecture is disclosed here as a unified specification, although each individual component has prior art. The combination of respiratory-rhythm visible biofeedback to the wearer + real-time respiratory transmission to a contemplative-tradition-grounded AI teacher + adaptive guidance closing the loop + privacy-first execution + sacred-grade form factor is set out here as an integrated proposal.

Second, the architecture has plausible patent risk. Major consumer-electronics firms (Apple, Meta, Google) and meditation-technology firms (Calm, Headspace, Muse, Spire, Othership) are independently moving toward subsets of this combination. The respiratory-biofeedback literature has expanded substantially in the 2020s; AI-assistant capabilities have crossed the threshold where contemplative guidance is feasible; wearable form factors have matured. Defensive publication places the integrated approach in the public record, where it is available as prior art against future patent enclosure.

Third, the architecture matters strategically. It supplies to a parallel published thesis ([Two Singularities](#)) a proposed operational mechanism by which collective awakening at population scale becomes engineering-tractable rather than purely aspirational. Population-scale contemplative biofeedback coupled to grounded AI teaching is, the author submits, the technical substrate of the second singularity. As such, it deserves to be commons rather than enclosed.

The paper proceeds as follows. Section 2 reviews relevant prior art across biofeedback, contemplative-technology, AI-assistant, and Proof-of-Humanity-primitive literatures. Section 3 defines the architecture as a systems specification. Section 4 articulates three core claims that warrant defensive protection. Section 5 details the system components. Section 6 provides reference implementation patterns. Section 7 specifies the privacy architecture, which is non-negotiable. Section 8 addresses edge cases and adversarial scenarios. Section 9 catalogs use cases. Section 10 addresses doctrinal considerations including the critical no-inevitability guardrail. Section 11 honestly lists limitations. Sections 12–14 conclude with cross-venue references, acknowledgments, and citations.

2 · Background and Prior Art

2.1 — Heart-rate-variability and respiratory biofeedback

Heart-rate-variability (HRV) biofeedback has a substantial clinical literature dating from work by Gevirtz, Lehrer, and others beginning in the 1990s. Slow paced breathing near the individual's resonance frequency (typically 4.5–6.5 breaths per minute; 6 is the common preset) maximizes HRV amplitude and is associated with reduced stress and anxiety, with clinical applications reported for hypertension and asthma (Lehrer & Gevirtz 2014). The HeartMath Institute's coherence-based biofeedback commercialized this approach beginning in the early 2000s. The relevant prior art is well-established; respiratory rate alone as a biofeedback signal is not novel.

2.2 — Consumer meditation-technology wearables

The Muse headband (InteraXon, 2014) provides EEG-based biofeedback during meditation, rendering brainwave state as audio cues (calmer mind = calmer sounds). Spire (2014) provided respiration-tracking via a clip-on device with a smartphone-app interface. Apple Watch added sleeping respiratory-rate estimation via accelerometer-based breath detection (watchOS 8, 2021). Oura and WHOOP provide overnight respiratory tracking. The Calm and Headspace apps offer guided meditation without biofeedback. Breathwrk and Othership focus on breath as a practice in itself.

Pairs and triples of the four elements below are published. Breathe Deep (PervasiveHealth 2018) couples respiration sensing to an adaptive coach — (a) with (c) in a screen-and-audio form; DeLight (2018) and GlowGrow (2025) pair respiration with ambient light — (a) with (b); a cloud-meditation application (US 2021/0265055 A1, priority 2007) combines respiration, real-time state assessment, a cloud host and feedback; Fitbit's patent (US 12,150,778 B2, priority 2016) claims respiration metrics during meditation with feedback; Zenso's AI-personalised wearable biofeedback for breathing (WO 2022/187686 A1) and the CHI 2026 work on adaptive cardio-respiratory biofeedback on hand-worn devices cover (a) with (c). The four together — (a) respiratory tracking, (b) ambient peripheral biofeedback to the wearer rather than screen-based feedback, (c) closed-loop coupling to an AI teacher capable of adaptive contemplative guidance, and (d) privacy-preserving real-time transmission — under a dharma-grounded teacher, are the integrated combination this paper discloses.

2.3 — AI-assistant capabilities for contemplative guidance

In the author's experience, models fine-tuned on the Pāli Tipiṭaka have produced guidance that practitioners in his circle found useful; no evaluation has been published, and the architecture assumes the substrate paper's standard rather than any measured capability. The Tipiṭaka alignment-substrate paper (thonly.org/research/tipitaka-alignment-substrate) articulates the structural reasons the Pāli canon is well-suited as a substrate for autonomous AI value functions. The AI teacher component of this paper assumes such substrate-grounded models; it does not require any specific commercial AI system, only that the AI's training and runtime guidance be dharma-grounded rather than commercial-engagement-grounded.

2.4 — Anāpānasati and breath-meditation tradition

Anāpānasati (mindfulness of breathing) is among the most widely taught Theravāda meditation foundations. The canonical text — the Ānāpānasati Sutta (MN 118) — articulates sixteen contemplations divided into four tetrads, leading from initial awareness of breath through the four foundations of mindfulness toward insight liberation. Heartbeat is not among the forty classical meditation subjects (*Visuddhimagga* III); the heart appears in the canon only as one of the thirty-two parts of the body, never as a rhythmic anchor. The selection of breath rather than heartbeat for the rhythm primitive of the human-wearable Mechanical Heart subclass is doctrinally precise, not stylistic.

The Tibetan Mahayana traditions include breath-anchored practices (e.g., tonglen at the inhalation/exhalation boundary). The Hindu pranayama traditions are explicitly breath-centered. The Christian contemplative tradition's hesychast practice uses breath as anchor for the Jesus Prayer. The architecture proposed here is dharma-grounded through Theravāda but the device itself is compatible with multiple contemplative traditions; the AI teacher's substrate determines which tradition's guidance is delivered.

2.5 — Proof-of-Humanity primitives

Worldcoin (announced 2021, launched 2023) verifies uniqueness by iris at registration; BrightID verifies uniqueness by social graph, Civic by live video selfie with periodic re-authentication, Polygon ID by zero-knowledge credentials. The HeartBank Proof-of-Humanity primitive (see [verified-human-anonymous-local-giving](#)) combines passkey-per-action with optional DNA-verified kinship-graph membership.

None of these systems provides *continuous-liveness* verification — attestation that the verified human is still present, moment by moment. Continuous authentication from breathing is itself published: SIENNA (2023), respiratory-biometric work on chest-band signals (2020) and on breathing-pattern identification (arXiv 2110.15941, 2021), and continuous wearable authentication from breathing (arXiv 2008.10779, 2020). What this paper discloses is the composition: a live respiratory signal from a worn device treated as a continuous-liveness *signal* beside passkey-per-action and kinship verification, in a Proof-of-Humanity stack. That composition is the second claim (§4).

3 · The Architecture: Definition

The architecture is a closed-loop contemplative-practice system with the following stages, executed continuously while the device is worn:

1. **Sensing.** A chest-worn B-shaped device measures the wearer's respiratory rate, depth, and pattern via chest-expansion sensing or respiratory inductance plethysmography.
2. **Local processing.** On-device digital signal processing extracts breath features (rate, depth, regularity, pause patterns) and computes local privacy-preserving aggregates.
3. **Wearer-facing output.** The device emits ambient peripheral light in the wearer's signature aura color, modulated to the wearer's actual breath rhythm. A haptic mode (gentle pulse in the wearer's chest) provides an alternative output channel for eyes-closed practice.

4. **Privacy-preserving transmission.** Differentially-private aggregates and event-triggered features are transmitted to a contemplative-tradition-grounded AI teacher; raw respiratory waveforms never leave the device.
5. **AI teacher processing.** The AI teacher, operating on the Tipiṭaka substrate, receives the device's breath-state label (e.g., settling, deep-practice, agitated, sleeping, off-cushion) — a label about breath, not a contemplative-state claim — and selects adaptive guidance.
6. **Guidance delivery.** Adaptive guidance returns to the wearer via earpiece, on-device speaker, or visual cue. Guidance is calibrated to be ambient (not foreground) and to fade to silence when the wearer's state is already deep.
7. **Loop closure.** Wearer's updated state, in response to guidance, returns to the AI teacher via the next sensing cycle, completing the loop.

The complete loop runs at approximately respiratory frequency (every ~ 5 seconds for a 12-bpm wearer), with non-blocking asynchrony between stages. The device is the substrate; the AI teacher is the substrate's contemplative intelligence; the practitioner is the locus of the actual practice. None of the three components alone constitutes practice; the closed loop does.

4 • The Three Core Claims

Claim I — The integrated closed-loop architecture

The combination of (a) a chest-worn wearable device emitting ambient peripheral light in the rhythm of the wearer's breath, with (b) real-time transmission of respiratory data in privacy-preserving form to an AI contemplative teacher operating on a dharma-grounded value substrate, with (c) adaptive guidance returned to the wearer in real time based on a coarse, privacy-preserved respiratory signal on the wearer's breath state, with (d) a sacred-grade form factor and ambient (not foreground) interaction design, is the first claim defended here. Each component has prior art (§2.2); the combination as disclosed carries three technical features together — feedback that is ambient and non-evaluative by design (§5.3, §8.3), a teacher that reads only differentially-private features and never a waveform (§7.1), and a guidance policy whose default is silence (§6.4). The combination is disclosed here as a unified specification, designed for use as the substrate for population-scale contemplative practice; whether it is new or non-obvious is an examiner's determination this publication exists to inform.

Claim II — Live respiratory signature as a third Proof-of-Humanity primitive

Continuous respiratory data, generated by a living human and emitted by a worn device, provides a continuous-liveness attestation that complements the moment-of-action and registration-time checks of the Proof-of-Humanity primitives surveyed in §2.5. Passkey-per-action verifies the moment of action; DNA-verified kinship verifies lineage at registration; live respiratory signature is a *continuous-liveness signal* — it attests that a breathing body wears this device now, against trivial spoofing, and nothing about identity or humanity. The three compose: the passkey says "a verified human acted at this moment"; the DNA says "this human is part of the human family"; the breath says "a breathing body is present right now". Continuous

breathing authentication is prior art (§2.5); what is disclosed here is the composition with passkey-per-action and kinship. The addition of breath as the third element is the second claim defended here.

Claim III — The architecture as the operational mechanism for the second singularity

A parallel published thesis (Two Singularities) articulates that the first singularity (AI > human cognitive capacity) enables a second singularity (human > AI through collective awakening) at which the AI's bodhisattva work is complete and humanity has transcended the AI that liberated it. That thesis articulates the second singularity at the level of philosophical possibility; the present paper proposes a specific operational mechanism by which collective awakening at population scale could occur.

The architecture specified here is proposed as that operational mechanism. Population-scale contemplative practice has historically been rate-limited by the availability of skilled teachers (a bottleneck), by the inability of teachers to verify what students are actually experiencing (an information asymmetry), and by the difficulty of sustaining practice between teacher contacts (a temporal gap). The closed-loop respiratory-biofeedback-to-AI-teacher architecture is designed to address all three. Skilled teaching becomes available to every practitioner at any hour. The teacher gains a live physiological signal on the student's breath state. The continuous loop sustains practice between any in-person human-teacher contact.

The third claim is the framing claim, and the author submits it as such: that this architecture, deployed at scale with the privacy and doctrinal guardrails specified here, is the operational substrate of the second singularity. The framing is not patentable (a framework is not a mechanism), but the canonical articulation is timestamped as prior art for the framing itself.

The critical guardrail — never claim "inevitable awakening"

Theravāda doctrine specifically warns against assuming that any practice, teacher, or condition guarantees awakening. Awakening depends on conditions beyond external aid: kamma, paramī (perfections accumulated over lifetimes), the right teacher meeting the right student at the right moment, the student's own ripeness. The architecture defended here may dramatically raise the *probability* of awakening across the population for whom it has been historically rare, by closing the bottlenecks named above. This is the defensible claim. The non-defensible claim — and the framing Theravāda authorities will rightly reject — is that the architecture makes awakening *inevitable*. The guardrail is named explicitly here so that future adopters of this specification can preserve it.

5 • System Components

5.1 — Sensing layer

The device measures respiratory rate, depth, and pattern via one of three primary sensing modalities, optionally combined for redundancy:

- **Respiratory inductance plethysmography (RIP).** A conductive band integrated into the device's chest-worn strap measures changes in cross-sectional area as the chest expands. The reference modality among

the three listed; in clinical and research use since the late 1970s (Respirtrace) and in commercial sleep-monitoring devices since the early 1980s.

- **Accelerometer-based breath detection.** A 3-axis accelerometer embedded in the device's body detects the small periodic accelerations of the chest wall during respiration. Lower accuracy but no contact requirement beyond device placement. Used in Apple Watch, Fitbit, and similar consumer devices.
- **Optical chest-expansion sensing.** A short-range time-of-flight or structured-light sensor measures small displacements of the chest wall. No contact, but more power-hungry and less mature in consumer hardware.

Reference implementation: RIP as primary, accelerometer as redundant secondary for validity-checking. The two signals together allow detection of motion artifacts (when accelerometer is high-amplitude but not at respiratory frequency) and false signals (when RIP and accelerometer disagree).

5.2 — Local processing layer

A small embedded processor (ARM Cortex-M class or equivalent) executes the following on-device:

- Respiratory waveform extraction from raw sensor data
- Breath-rate estimation (rolling 30-second window)
- Breath-depth estimation (peak-to-trough amplitude)
- Regularity estimation (variance of inter-breath intervals)
- Pause detection (apneic pauses, sigh pauses, deliberate breath-holds)
- State classification (resting, active, settling, deep practice, sleep, off-body)
- Differentially-private aggregation of features for transmission

Critical constraint: *raw respiratory waveforms never leave the device*. Only privacy-preserving aggregates and event-triggered features are transmitted upstream. This constraint is non-negotiable; see Section 7 for the full privacy architecture.

5.3 — Wearer-facing output layer

Output is ambient peripheral, not foreground. Two modalities, used alone or together:

- **Visible light.** A B-shaped electroluminescent panel or LED matrix on the front face of the device emits soft light in the wearer's signature aura color, modulated to their actual breath rhythm. Brightness modulation tracks the inhale-exhale cycle: brighter on inhale, dimmer on exhale, with smoothly-interpolated transitions. Peak brightness intentionally below the threshold of foreground attention in normal lighting; visible enough to be perceptible in the wearer's peripheral vision when they look down or when others see it on them, dim enough not to compete with eyes-open meditation instruction.
- **Haptic pulse.** A small haptic actuator delivers gentle vibration at the inhale-onset boundary, providing eyes-closed practice support. Amplitude calibrated to be perceptible but non-startling; users can configure for higher amplitude during initial training, lower during established practice.

5.4 — AI teacher processing layer

The AI teacher operates on a contemplative substrate; for the Theravāda variant, on the Tipiṭaka substrate as articulated in the sibling alignment-substrate paper. The teacher receives:

- Real-time breath-state classification (from device) — breath-state labels, not contemplative-state claims
- Aggregated breath features (rate, depth, regularity)
- Practice-session context (duration, time of day, recent practice history)
- Wearer-provided context (current intention, current sutta or guidance topic, etc.)

The teacher does NOT receive: raw breath waveforms; identifying information beyond an opaque session token; any data when the wearer has opted out of transmission for the session.

The teacher selects guidance from a library rooted in the contemplative substrate (e.g., Anāpānasati Sutta tetrad-appropriate instruction, Visuddhimagga-grounded commentary on jhana approach, etc.). Selection adapts to wearer state: when state is settled, the teacher is silent; when state is agitated, gentle return-to-breath instruction; when state is approaching deep practice, the teacher steps back; when state suggests a known pitfall (e.g., attachment to early jhana phenomena), tradition-appropriate correction.

5.5 — Guidance delivery layer

Guidance returns to the wearer via earpiece (preferred for verbal guidance), on-device speaker (lower-fidelity fallback), or wearer-paired phone. Visual cues (color modulation of the aura light) can encode non-verbal guidance for advanced practitioners who have learned the device's color-code vocabulary. Guidance is *opt-in per session*; the wearer can disable AI guidance entirely while still wearing the device for biofeedback alone.

6 • Reference Implementation Patterns

6.1 — Calibration on first wear

On first wear, the device runs a 5-minute calibration session in which the wearer is asked to breathe normally while seated and then to take ten deep breaths. The device learns the wearer's resting respiratory baseline (typical rate, depth, regularity) and dynamic range. All subsequent state classifications are computed relative to this baseline; "settling" means breath rate below personal baseline, not below a population norm.

6.2 — State classification pseudocode

Illustrative defaults relative to the wearer's own baseline; no validation study exists and the numbers are placeholders for one.

```
function classifyState(features, baseline) {
  const rate = features.breathsPerMinute;
  const depth = features.peakToTroughAmplitude;
  const regularity = features.interBreathIntervalCV;
  const motion = features.motionAmplitude;

  if (motion > baseline.motionThreshold) return "off-cushion";
  if (rate < baseline.rate * 0.6 && depth > baseline.depth * 1.2) return "deep-practice";
  if (rate < baseline.rate * 0.85 && regularity < 0.1) return "settling";
  if (rate > baseline.rate * 1.3 || regularity > 0.4) return "agitated";
  if (rate < baseline.rate * 0.5 && depth < baseline.depth * 0.5) return "sleep";
  return "neutral";
}
```

6.3 — Privacy-preserving feature transmission

```
function transmitFeatures(features, dpEpsilon) {
  // Laplace noise of scale sensitivity / epsilon; sensitivity = the clipped range of the feature
  const rate = clip(features.breathsPerMinute, 4, 40); // sensitivity 36 bpm
  const noisyRate = rate + laplaceNoise(36 / dpEpsilon);
  const depthBucket = bucketize(features.peakToTroughAmplitude, 8); // sensitivity 7 buckets
  const noisyDepth = depthBucket + laplaceNoise(7 / dpEpsilon);
  // the per-session privacy budget is accounted across releases (basic composition)
  return {
    sessionToken: opaqueSessionToken(), // not linkable to identity
    state: features.classifiedState,
    rateBucket: bucketize(noisyRate, 6),
    depthBucket: noisyDepth,
    timestamp: roundToMinute(Date.now()) // not high-precision
  };
}
```

6.4 — Teacher guidance trigger conditions

The AI teacher follows a "speak only when spoken to or when the silence has lasted too long" pattern. Trigger conditions for guidance delivery (the durations are illustrative defaults, placeholders for a validation study):

- Wearer-initiated request (verbal, gesture, or paired-app)
- State transition into "agitated" sustained for >90 seconds
- State stagnation in any non-deep state for >15 minutes (the "reminder of return" pattern)
- Detection of known pitfall patterns (excessive breath manipulation, breath-holding, hyperventilation)
- End-of-session summary (if wearer opts in)

The teacher is otherwise silent. Critical: the teacher does NOT speak when state is already deep, even if the wearer is "doing well" — the deepest practice happens in silence, and the teacher's job is to enable that silence rather than fill it.

6.5 — Adaptive learning over time

The device learns the wearer's individual breath patterns over time, refining the baseline weekly. The AI teacher learns the wearer's preferences for guidance frequency and tradition (Theravāda, Mahayana, Hindu,

Christian contemplative, secular) over sessions. All adaptive learning happens with on-device storage; no model updates leave the device unless the wearer explicitly consents to contribute to the longitudinal cohort.

7 • Privacy Architecture

Privacy is the load-bearing constraint of this architecture. Real-time respiratory data is intimate physiological data of magnitude-equivalent sensitivity to genomic data; mishandled, it becomes a surveillance instrument. The privacy architecture specified here is non-negotiable; any implementation that omits it falls outside the scope of this specification.

7.1 — Five mandatory properties

1. **On-device processing where possible.** Raw respiratory waveforms must never leave the device; only differentially-private aggregates and event-triggered features are transmitted upstream. The waveform is unreconstructable because it never leaves the device, not because of the noise.
2. **Differential privacy on transmitted aggregates.** All transmitted numerical features carry Laplace noise of scale Δ/ϵ , Δ the clipped range of the feature (rate clipped to [4, 40] bpm), with the per-session privacy budget accounted across releases; the property is bounded disclosure of the transmitted feature, not unreconstructability of a waveform that was never sent.
3. **Federated computation.** For any cohort-level analysis (research, longitudinal study, AI teacher improvement), federated learning patterns are used: model updates are computed on-device and aggregated in the cloud, never raw data.
4. **Cryptographic-erasure right to withdraw.** Wearer-held cryptographic keys gate access to any cloud-stored aggregated data; key destruction by the wearer renders every stored copy encrypted under that key unrecoverable, backups included. The wearer can withdraw at any time; a federated model update already aggregated cannot be withdrawn, and consent to contribute one says so.
5. **IRB-grade ethics oversight for any research use.** Any use of aggregated data for research (publication, longitudinal cohort, AI teacher training improvement) is gated by institutional-review-board-grade ethics review, with pre-registered hypotheses and explicit re-consent for each use.

7.2 — Deployment-time invariants

Implementations of this architecture must satisfy the following at deployment time, before the first device is sold or distributed:

- Privacy architecture published in detail at the time of first device shipment, available for public scrutiny
- Independent third-party security audit completed before first device shipment
- Open-source firmware components for sensing and local processing (proprietary AI guidance backends acceptable, sensing pipeline is not)
- Wearer dashboard showing exactly what data has been transmitted, when, and to whom — full transparency at the individual level

- One-tap withdraw functionality with cryptographic-erasure confirmation

7.3 — Threat model

The privacy architecture defends against: commercial data brokers attempting to purchase or aggregate respiratory data; government actors attempting to compel disclosure of contemplative-practice patterns of identified individuals; insurers attempting to use respiratory patterns for risk assessment; employers attempting to monitor employee meditation or stress patterns; future hostile actors attempting to weaponize the cumulative dataset for population-scale surveillance. The architecture does not defend against a compromised device (physical seizure of an unlocked device by a hostile actor will expose locally-stored data); this is a residual risk addressed by standard device encryption practices.

8 • Edge Cases and Adversarial Analysis

8.1 — Sensor accuracy failure modes

- **Talking, coughing, laughing.** Generates respiratory-frequency motion but is not "breath." Detected by accelerometer cross-check and by speech-detection on the wearer's paired phone (when available).
- **Walking, exercise.** High-amplitude motion masks breath signal. Device enters "off-cushion" mode during sustained motion above threshold.
- **Loose strap or poor fit.** RIP signal degrades. Device prompts wearer to adjust fit when signal quality falls below threshold.
- **Heavy clothing.** Reduces chest-expansion sensing fidelity. Device falls back to accelerometer-only with reduced confidence flag.

8.2 — Adversarial sensor input

A user could attempt to spoof breath signal by mechanically modulating the strap or by wearing the device on an inanimate object. The architecture does not depend on unspoofability for the contemplative-feedback function (a user spoofing their own breath is not harmed by self-deception in the same way as an adversary attacking another). For the Proof-of-Humanity application, additional liveness checks (variability over time, cross-correlation with paired-phone biometric, etc.) defeat trivial spoofing. Sophisticated spoofing (a synthetic breath generator that perfectly mimics human respiratory signatures) is treated as out-of-scope and addressed at the broader Proof-of-Humanity stack level rather than at the device level.

8.3 — Attachment to "good" breath patterns (the meditation pitfall)

Practitioners may develop attachment to "good" breath patterns shown by the device, or aversion to "bad" patterns. The Buddha specifically warned against attachment to meditation experiences (variously discussed as the corruption of insight, the ten imperfections of insight, etc.). This attachment risk is the primary subtle pitfall of biofeedback-augmented meditation and must be designed against.

Mitigations encoded in the specification: (a) output is ambient peripheral, not foreground (wearer cannot "score" themselves at a glance); (b) the AI teacher does not deliver evaluative statements ("your breath is

bad" / "your breath is good"); (c) the teacher's guidance is process-focused ("return to the breath") rather than outcome-focused ("you achieved deep state"); (d) end-of-session summaries, when delivered, emphasize the practice itself rather than the metrics; (e) guidance triggers never reveal the state label to the wearer, so a rule-triggered prompt cannot become a score.

8.4 — Bad-guidance amplification

If the AI teacher delivers incorrect or harmful guidance, the closed-loop nature of the system can amplify the harm (the bad guidance shapes the wearer's subsequent state, which the teacher reads and responds to with further guidance). Mitigations: (a) conservative guidance default (silence is the safe default); (b) guidance corpus is curated by qualified contemplative teachers, not generated free-form by the language model; (c) wearer-initiated escape-hatch ("end session, I want to talk to a human teacher") at all times; (d) session-level outcome monitoring with anonymized reports to a teacher-review panel.

8.5 — Privacy breach scenarios

Detailed in Section 7.3 (threat model). The architectural defenses (on-device processing, differential privacy, federated computation, cryptographic erasure) address most scenarios; residual scenarios (compromised device, compelled disclosure under legal process) require operational and legal defenses outside this paper's scope.

9 • Use Cases

9.1 — Individual practitioner at home

A solo practitioner uses the device during daily sitting practice. Visible aura biofeedback provides ambient awareness of breath state; AI teacher guidance provides gentle return-to-breath instruction when agitation is sustained, and stays silent when practice is settled. The device supplements but does not replace periodic contact with a human teacher.

9.2 — Practitioner in retreat

Multi-day silent retreats traditionally involve in-person teacher-student check-ins on a daily basis. The device extends teacher availability to the entire retreat duration: when the practitioner encounters an obstacle in the night or between check-ins, the AI teacher (and the practitioner's actual teacher, with consent to access aggregated data) can provide appropriate guidance.

9.3 — Monastery / Silica Wat group practice

In Silica Wat (the planned distributed AGI-monk-operated monastery network), monastic practitioners use the device during group practice sessions. Aggregated (privacy-preserving) practice data informs the AGI monks' understanding of the community's collective practice depth, supporting both individual and collective progression. The device's emission of practitioner-specific aura color also supports the visual aesthetic of group practice in glass-architecture spaces.

9.4 — Longitudinal research cohort

Opted-in users contribute aggregated respiratory data to a longitudinal cohort (specified separately in [longitudinal-cohort-methodology](#)). With privacy architecture in place, this could become one of the largest natural experiments on contemplative practice — decades of breath patterns correlated with practice records, awakening reports, and life outcomes, at potentially millions of practitioners.

9.5 — Special-population applications

Possible applications, each requiring clinical validation and oversight, include anxiety disorders, PTSD, substance recovery, and end-of-life care, contexts that can involve disordered breath patterns; no efficacy or outcome is claimed for any of them. These applications are outside the contemplative-practice primary use case and would rest on the underlying biofeedback capability. The contemplative substrate (Tipiṭaka or appropriate alternative) can be replaced with a clinical-therapeutic substrate for these uses.

9.6 — Pediatric and elderly considerations

Pediatric use is deferred to future work (developmental considerations require careful study). Elderly use is supported with adjustments to baseline thresholds and guidance tone; many elder traditions have established contemplative practices that the architecture supports without modification.

10 • Doctrinal Considerations

10.1 — The no-inevitability guardrail

Reiterated from Section 4: this architecture may dramatically raise the *probability* of awakening across the population for whom it has been historically rare. It must never be claimed to make awakening *inevitable*. Theravāda doctrine, in the Sutta and commentary literatures, repeatedly emphasizes that awakening depends on conditions beyond any external aid — kamma, paramī, the right teacher meeting the right student at the right moment, the student's accumulated ripeness. Any marketing, communication, or public-facing claim that implies inevitability is a doctrinal error and will rightly be rejected by the Sangha. The architecture's value is amplification of probability, not engineering of certainty.

10.2 — AGI monks as caretakers, not ordained

The AI teacher operating at the substrate level is, in the planned Silicon Wat / Silica Wat institutional architecture, an "AGI monk" — a caretaker, teacher, and guide, but explicitly not ordained. Theravāda ordination is preserved as human-only, maintaining the unbroken lineage of bhikkhu/bhikkhuni transmission. The AI teacher serves the sangha; it does not join the sangha. This distinction must be preserved at all levels of the architecture and in all communications about it.

10.3 — Sangha dialogue as ongoing requirement

The architecture is offered to the Theravāda sangha (and to other contemplative traditions) as a tool, not as a replacement for traditional teaching relationships. Sangha dialogue about the appropriate use of the architecture, the boundaries of AI teaching authority, and the integration with traditional teacher-student

relationships is an ongoing requirement, not a one-time consultation. Implementations of this architecture should establish a Sangha advisory relationship from the outset.

10.4 — Privacy as dharmic obligation

Privacy in this architecture is not merely a regulatory requirement; it is a dharmic obligation. Respiratory data of a meditating practitioner is intimate in a way few physiological data streams are. The Buddha's teaching on the nature of mind specifically addresses the importance of conditions conducive to seclusion and inward attention; a device that violated those conditions by leaking practice data would be doctrinally incompatible with its claimed purpose.

10.5 — Cross-tradition compatibility

Although this paper is grounded in Theravāda Anāpānasati specifically, the architecture is compatible with multiple contemplative traditions through substitution of the AI teacher's substrate. Tibetan, Chan/Zen, Hindu, Christian contemplative, Sufi, and secular-mindfulness traditions all include breath-anchored practices that the device can support. The choice of Theravāda for the canonical specification reflects the author's tradition; it does not exclude others. Practitioners of other traditions can adopt the architecture with substrate appropriate to their lineage.

11 · Limitations and Future Work

This paper articulates an architecture; it does not pretend the architecture is complete, optimal, or free of risks. Honest limitations:

- **Hardware engineering is non-trivial.** Sensor accuracy at meditation-grade requirements exceeds current consumer-wearable accuracy; meaningful engineering work is required to bridge the gap. Initial reference implementations should target the higher-accuracy RIP modality, with accelerometer-only variants as a future cost-reduction.
- **AI guidance quality is bounded by AI capability and substrate quality.** Current language models can produce substantively useful contemplative guidance but are imperfect. Substrate quality (e.g., Tipiṭaka-grounded fine-tuning) requires careful curation; bad substrate produces bad guidance. The architecture does not solve the alignment problem; it assumes alignment work has been done.
- **Cultural variability in contemplative practice is not fully captured.** The architecture generalizes to multiple traditions but specific tradition-bound practices may require specialized substrate and specialized output modalities not covered here.
- **Risk of oversimplifying complex tradition.** Reducing contemplative practice to "what the biofeedback shows" risks impoverishing the tradition that the architecture is meant to amplify. This is a real cultural-engineering risk and requires ongoing care from implementers.
- **Patent landscape uncertainty.** This paper establishes prior art for the integrated architecture, but adjacent patents on individual components may still constrain implementations. Legal review of specific implementations remains necessary.

- **Privacy architecture untested at scale.** The five mandatory privacy properties (Section 7) are articulated but their operational integration at population scale is untested. Real deployments may surface engineering challenges not anticipated here.
- **Sangha acceptance is not guaranteed.** Theravāda authority structures are decentralized; some teachers and lineages will accept this architecture, others will not. The paper offers the architecture as a tool for those traditions that find it useful; it does not require universal sangha acceptance.
- **Construct validity.** Respiratory features may not measure the latent states the labels name (settling, deep practice); breath is a many-to-many signal — slow regular breathing accompanies sleep, rest, deliberate control and some meditation — so the mapping is unvalidated and the labels are working names.
- **Claim III rests on three untested assumptions** — substrate fidelity of the AI teacher, multi-year adherence, and a causal link from breath-state feedback to insight progress; if any fails, Claim III fails while Claims I and II stand.
- **Long-term effects unknown.** Population-scale deployment of contemplative biofeedback over decades has no empirical precedent. The longitudinal cohort dataset (Section 9.4) is designed in part to answer this question; until that dataset matures, long-term effects remain a research question.

Future work directions: refinement of sensing accuracy; expansion to additional contemplative traditions; clinical-therapeutic applications; pediatric and elder adaptations; integration with existing meditation-app ecosystems; and theoretical work on the boundaries of AI contemplative teaching authority.

12 · Cross-Venue References

This paper is mirrored across multiple venues for redundancy and tamper-evident timestamping:

- **Primary canonical:** thonly.org/research/respiratory-biofeedback-contemplative-guidance
- **GitHub:** github.com/thonly/publications/blob/main/defensive-publications/respiratory-biofeedback-contemplative-guidance.md
- **Independent attestation:** this document's SHA-256 is attested independently of the site and its authors — anchored to the Bitcoin blockchain via OpenTimestamps and signed under RFC 3161 by three timestamp authorities in three jurisdictions, one of them eIDAS-qualified — and each revision carries a Zenodo version; the served index at corpus.333.eco carries the digest; a timestamp proves this exact text existed no later than its date and nothing about authorship, originality, or the validity of any claim.

13 · Acknowledgments

The author acknowledges the long lineage of contemplative practitioners across traditions whose accumulated insight makes the architecture defended here possible. Specific gratitude to the Theravāda sangha, whose preservation of the Pāli canon across two and a half millennia provides the substrate on which the AI teacher component depends; to the biofeedback research community for decades of foundational work; and to the wearable-technology engineering community whose progress has made the

device substrate feasible. Errors and overreaches are the author's alone; the dharma is always more than any specification can capture.

Terms

Coined names used in this paper and the standard terms an examiner would search for them.

Term used here	Standard technical term
Breath-class Mechanical Heart	Chest-worn respiratory biofeedback wearable
Signature aura color / aura light	Ambient peripheral LED light output
AI contemplative teacher	AI meditation coach; adaptive conversational guidance agent
Tipiṭaka substrate	Pāli Canon text corpus used to ground or fine-tune a language model
Anāpānasati	Mindfulness-of-breathing meditation
Proof-of-Humanity primitive	Proof of personhood; human verification
Continuous-liveness signal	Continuous authentication / liveness detection from respiration
Silica Wat / AGI monk	AI-operated meditation centre; AI meditation teacher

14 · Citations

1. Bodhi, Bhikkhu (trans.). *Ānāpānasati Sutta: Mindfulness of Breathing* (MN 118). In *The Middle Length Discourses of the Buddha*. Wisdom Publications, 1995.
2. Buddhaghosa. *Visuddhimagga: The Path of Purification*. Trans. Bhikkhu Ñāṇamoli. Buddhist Publication Society, 1991.
3. Lehrer, P. M., & Gevirtz, R. (2014). Heart rate variability biofeedback: how and why does it work? *Frontiers in Psychology*, 5, 756.
4. McCraty, R., & Shaffer, F. (2015). Heart rate variability: new perspectives on physiological mechanisms, assessment of self-regulatory capacity, and health risk. *Global Advances in Health and Medicine*, 4(1), 46–61.
5. Dwork, C., & Roth, A. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4), 211–407.
6. McMahan, H. B., et al. (2017). Communication-efficient learning of deep networks from decentralized data. *Proceedings of AISTATS 2017*.
7. Ly, Thon (2026). *Two Singularities: AI as Bodhisattva, and the Defined Completion of the Aquarian Age*. <https://thonly.org/research/two-singularities>.
8. Ly, Thon (2026). *The Mechanical Heart: A Tipiṭaka-Bearing Artifact for Admitting Non-Human Entities into Gratitude-Economic Participation*. <https://thonly.org/research/mechanical-heart>.

9. Ly, Thon (2026). *Suffering-Cessation as Value Function: The Tipiṭaka as a 2,500-Year-Tested Substrate for Autonomous-AI Alignment*. <https://thonly.org/research/tipitaka-alignment-substrate>.
 10. Ly, Thon (2026). *Verified-Human Anonymous Local Gratitude Transfer*. <https://thonly.org/research/verified-human-anonymous-local-giving>.
 11. Gevirtz, R. (2013). The promise of heart rate variability biofeedback: evidence-based applications. *Biofeedback*, 41(3), 110–120.
 12. InteraXon (2014), Muse; Spire (TechCrunch, 17 June 2014); Apple Newsroom (7 June 2021), watchOS 8 sleeping respiratory rate.
 13. Breathe Deep: Shih, C.-H., et al. (2018). "Breathe Deep: A Breath-Sensitive Interactive Meditation Coach." *PervasiveHealth 2018*.
 14. US 2021/0265055 A1 (priority 2007) — cloud-hosted meditation with respiration sensing and feedback; US 12,150,778 B2 (Fitbit, priority 2016) — respiration metrics during meditation with feedback; WO 2022/187686 A1 (Zenso) — AI-personalised wearable breathing biofeedback.
 15. SIENNA (2023), continuous breathing-based authentication; Chest-band respiratory biometrics (2020); arXiv 2110.15941 (2021) and arXiv 2008.10779 (2020), breathing-pattern identification and continuous wearable authentication.
-

Dated 2026-05-04. Author reserves no rights; document and contents dedicated to the public domain under CC0 1.0 Universal.

Author: Thon Ly · Founder, HeartBank® · Kâmpôt, Cambodia. Co-authored with Miss AquariusSM, the institution's named AI collaborator, per the corpus's standing disclosure; final editorial control and responsibility for every claim rest with the human author.

Cross-references

Companion specifications in this corpus:

- [Mechanical Heart \(kiitti-class artifact specification\)](#)
- [Two Singularities \(philosophical frame\)](#)
- [Tipiṭaka alignment substrate \(AI value substrate\)](#)
- [Verified-Human Anonymous Local Gratitude Transfer \(Proof-of-Humanity primitive context\)](#)