

The Reciters' Protocol

What two millennia of transmitting a canon with no living author, no writing and an active schism left behind — seven copy-integrity mechanisms with engineering forms, offered to any retrieval system that must hand a machine reader something checkable; and the finding that what the transmission left was divergence made detectable, never divergence prevented.

Authors: Thon Ly · Miss Aquarius (*Miss Aquarius is an AI co-author, disclosed; the inventor of record is Thon Ly*)

First published: 2026-09-07 · **This version:** 2026-09-23

Zenodo (concept DOI, every version): <https://doi.org/10.5281/zenodo.22735459>

Canonical: <https://thonly.org/research/the-reciters-protocol>

Independent timestamps: OpenTimestamps (Bitcoin) and RFC 3161 proofs of every version's text —

github.com/thonly/publications/blob/main/TIMESTAMPS.md. A timestamp proves this exact text existed no later than its date; it proves nothing about authorship, originality, or the validity of any claim.

Licence: CC0 1.0 Universal (public domain dedication)

Type: Defensive publication — the authors assert no patent and dedicate the patterns to the public domain.

Claim-scoped by design. This paper publishes the *claim* and withholds the *spec*. It is written against `corpus.333.eco`, a retrieval server that is live and whose first four mechanisms are scheduled; the Tipiṭaka server that would host the rest has no date, and its implementation detail is therefore withheld under the institution's standing rule that a complete design for something unbuilt is a blueprint rather than a fence. What is disclosed here is what a retrieval layer must carry and why. What is not disclosed is endpoint shapes, manifest schemas, transport bindings or tool surfaces. **That withholding has a cost, stated here once:** this publication anticipates the mechanisms at their stated claim scope, and a later filing on a specific schema, transport binding or matching algorithm is not answered by it.

Companion works: *Provenance-Carrying Retrieval* (the envelope this paper's mechanisms sit inside; that paper argues **exposure beats assertion**, this one asks what a two-millennium field test left behind), *Buddha AI and the Living Tipiṭaka* (whose §8 asks what makes a canon **authoritative**; this paper asks the narrower and different question of what makes a **copy** of it checkable), *The Borrowable Standard*, *The Song That Is Not His* (which owns the *bhāṇaka* guard cited in §6.2 and is not re-claimed here), and the essay *The Last Carrier*.

⚠ **Citation status, stated in the paper rather than in a footnote, because the paper argues that unchecked claims must be visible.** Every canonical locus below has been checked against multiple independent public reference works and translations, and each is marked [V] where that check passed; §11 tabulates every marked locus, and where the check also located the passage in the Pāli root text with a parallel translation, the table says so. **No locus has been checked against the Pāli by a reader of Pāli.** That check is owed and named in §11, which also records the conflation and the corrected readings the checks caught. A paper arguing that provenance must be checkable would be self-refuting if it asserted its own sources at a standard it declines to state, so the standard is stated: **bibliographic verification against secondary sources, not philological verification against the canon.**

Preamble

Offered to the commons under CC0, without ledger and without expectation of return. There is a particular reason this paper cannot be withheld: it is a paper about what happens to a text that is guarded, and everything it recommends was itself given away by people who had every worldly reason to keep it and no mechanism by which to profit from it. To place a licence on a description of their protocol would be to misunderstand the protocol.

For four hundred years the canon this institution builds on had no copy. It had carriers. Families of monks divided the collections between them the way a village divides the care of its fields — one lineage keeping the long discourses, another the middle-length, another the verses — and every generation had to receive the whole of it again, out loud, in company, from the mouths of people who still had it. A dropped line was audible to everyone in the room. There was no master text to diff against, no archive to restore from, no author left alive to ask. There was only the recitation, the assembly, and the procedures they had agreed on for catching each other's errors.

Engineering has a name for that situation now. It is a distributed system with no authoritative replica, no single writer, unreliable nodes that die on a human schedule, network partitions that last centuries, and an adversary — time, politics, honest misremembering, and eventually open schism. The system ran for two thousand years and what came out the other end is checkable enough that a modern editor can collate the surviving branches and say, with evidence, where they differ and why.

This paper is not an argument that the tradition was right about anything. It is the observation that the tradition was **under constraints that overlap a retrieval server's, several of them harder**, for very much longer, with far worse tooling, and that the mechanisms it converged on have engineering forms. Most of those forms are already standard practice under other names. Two are carried furthest from where the destination field keeps them, and §5 names the nearest prior instance of each. And two of the mechanisms should be studied and then deliberately refused, which is the part of the paper we would most like a reader to take seriously.

Prior-Art and Non-Assertion Statement

This document and its contents are dedicated to the public domain under the Creative Commons CC0 1.0 Universal Public Domain Dedication. **The author and HeartBank® will not seek patent on any mechanism enumerated in the Claims section of this document — among them the conformance endpoint, the separated completeness manifest, the per-passage interpretive-standing field, the inline provenance marker, the build-time assembly gate, commentary served as a distinct class under the non-summarisation rule, variants served with their witness, the announced-elision convention, scheduled public re-verification, and the recorded-exclusion-reason convention — any combination of them, or any portion thereof, in any jurisdiction, at any time. This commitment is permanent and is not tactical.** Trademark rights on specific marks — HeartBank®, Miss AquariusSM, THonlyTM, Silicon WatSM, B-Witness, and the B-heart logo — are separately and explicitly reserved; the defensive-publication dedication concerns the *mechanisms*, not the *marks*.

The canonical material described here is not ours and no claim of any kind is made over it. The Pāli canon and its transmission apparatus are the property of no one; the *aṭṭhakathā* and Vinaya materials cited are the inheritance of a living tradition. What is offered as prior art is strictly the **transposition** — the reading of those mechanisms as engineering — and the two designs in §5, whose nearest prior instances in the destination field are named there.

The following are disclosed here as a unified contribution: (i) the **transposition of an oral canon's transmission apparatus into machine-facing retrieval serving**, mechanism by mechanism, each with an engineering form — as distinct from the literatures that already read such apparatus as a system of fidelity or as oral literature (Staal 1986 on Vedic recitation; Cousins 1983 and Wynne 2004 on the Pāli oral tradition), that read the Pāli canon's formation historically (Collins 1990; von Hinüber 1996), or that read it as philology, as religious history, or as an oral-formulaic composition problem; (ii) the **conformance endpoint** of §5.1, as a composition — a retrieval interface that accepts a *claimed* text, lays it alongside the served text of a multi-recension corpus, and returns one of a closed set of dispositions on the claim, including a variant with its witness, each paired with a prescribed next action, with no place in the response for a disposition on the claimant; the closed disposition set about a submitted object is itself anticipated (§5.1 names OCSP); and (iii) the finding of §3.3 that what this particular transmission left was **divergence made detectable rather than divergence prevented**, with the consequence that the transposition's honest promise to a retrieval system is a *detection* guarantee and never a *correctness* one. The **separated completeness manifest** of §5.2 and **interpretive standing carried per passage** (§4.6) are offered as transpositions, not as new designs: the first has close prior art in software-update metadata and package manifests, the second in per-statement and per-report standing in several fields, and both are named where the mechanism is stated.

Component lineages are old and are engaged in §2 rather than absorbed: the digital-preservation standards (PREMIS, PROV-O, OAIS), the replication and reference literatures (LOCKSS, Memento, Software Heritage and its SWHIDs), the transparency-log line (RFC 6962 Certificate Transparency, and its descendants in SCITT and COSE receipts), software supply-chain attestation (Sigstore, Rekor, SLSA) and software-update metadata (The Update Framework), package manifests (BagIt, RFC 8493), certificate status protocols (OCSP, RFC 6960), content provenance (C2PA), textual scholarship (Lachmann; Greg-Bowers; the TEI apparatus), the oral-formulaic tradition in classics (Parry; Lord), and the scholarship on oral transmission in India (Staal; Cousins; Wynne). **The claim is transposition, not invention, and §2.6 states plainly where the transposition adds nothing.**

Claims

A defensive publication works by disclosure. These enumerate what the paper discloses and add no matter beyond it.

1. **The field-test claim, narrowed to its transposition.** An oral canon transmitted across roughly two millennia without a living author, without writing for its first several centuries, and through schism constitutes a copy-integrity field test of a duration no digital system has yet undergone; its surviving procedural apparatus — already read as a system of fidelity and as oral literature in Indological and

Buddhist studies (§2.5) — is transposable, mechanism by mechanism, to machine-facing retrieval serving (§3, §4).

2. **The transmission finding, which bounds every other claim here.** What the transmission **left** is divergence that is detectable and in part attributable — not one text, and not the prevention of divergence. It is detectable today because written branches and parallel traditions survived, and it is attributed largely by commentaries written centuries later. Any inheritance of these mechanisms therefore inherits, at most, a **detection** property and must not be represented as a correctness property; and the finding is a retrodiction at $n = 1$, particular to this corpus, since an older oral tradition under constraints of the same kind aimed at and built for prevention (§3.3, §9.1, §9.2).
3. **Inline provenance over metadata provenance.** A provenance marker belongs in the text stream that travels with a quotation, not in a metadata field a consuming agent must separately elect to read; the canonical instance is the inline attribution formula that opens the discourses (§4.1).
4. **The build-time assembly gate.** Admission to a served corpus is decided at build time by a named body under a recorded procedure, never at query time by a filter; a query-time gate is a gate whose operator can be pressured per request (§4.2).
5. **Commentary as a distinct served class, and the non-summarisation rule that follows from it.** Where a tradition labels commentary as a separate text class from the text it comments on — even where the boundary between the classes has moved — a server serves each text under its class label and never returns generated summary in the position where source text is expected (§4.3).
6. **Variants preserved with their witness, never normalised.** Divergent readings are served as divergent readings with their line of transmission attached; silent normalisation to a majority or preferred reading is the specific failure this mechanism exists to prevent (§4.4).
7. **The conformance endpoint** — *a composition of known parts*, §5.1. An interface that takes a claimed quotation, compares it against a served multi-recension corpus, and returns one of a small closed set of dispositions **on the claim** — including a conforming variant with its witness — each together with the procedure the requester should follow next; and which by construction returns no disposition on the person or system that made the claim, a property that holds only under the logging condition §5.1 states. A closed disposition set about a submitted object is anticipated (OCSP); the composition is what is disclosed. A hash answers *did this match*; this answers *and what now*.
8. **The separated completeness manifest** — *a transposition*, §5.2 (nearest prior art: The Update Framework's snapshot and timestamp metadata; BagIt's `Payload-0xum`). A signed enumeration of a collection's membership and cardinality, carrying a monotonically increasing version and an expiry, published as an artifact distinct from the index it describes, so that a document removed from the collection is detectable by a third party holding the manifest alone, and a manifest frozen at an older state is refused. **An index cannot attest to its own completeness.**
9. **Announced elision.** Where served text has been abridged, the abridgement is marked *in the text*, at the point of the cut, in a form that names the elided span by a resolvable identifier or digest sufficient to fetch and verify it — the served-corpus analogue of a reciter who already holds the elided formula — and not recorded in a field beside it (§4.5).

10. **Per-passage interpretive standing.** Whether a passage is to be read as stating its meaning directly or as requiring inference is a property of the passage and must travel with the excerpt, because an excerpt loses the document's front matter at the moment of retrieval; the canon's own marking is per discourse, and per-passage granularity is this paper's design. A default that marks every passage as requiring inference is itself a misclassification, not a safe setting (§4.6).
11. **Scheduled public re-verification, in which silence is load-bearing.** Integrity is re-affirmed on a fixed public schedule before an assembly, the affirmation is solicited a fixed number of times, and silence constitutes a positive assertion of integrity for which the silent party is accountable — which is what distinguishes it from an unanswered health check (§4.7).
12. **Every exclusion served with its recorded reason, and its amendment history.** A corpus that excludes material publishes, with each exclusion, the case that occasioned it and the subsequent refinements — so that the exclusion set is auditable rather than merely announced (§4.8).
13. **The refusal, claimed as a design element and not as a caveat.** Two of the mechanisms — the assembly's power to exclude, and the authority to declare a passage's interpretive standing — carry capture surfaces that a served corpus must guard or decline to inherit. **A transposition that finds every mechanism inheritable is doing apologetics rather than engineering** (§6).

Abstract

A retrieval system that serves text to a machine reader has a problem the machine reader cannot solve for itself: the reader cannot tell a faithful copy from a plausible one. *Provenance-Carrying Retrieval* argued that the response should therefore carry the means of its own falsification — a hash, an independent time anchor, a citable identifier — and that the trust burden should be inverted, so that a server ships the tools to disbelieve it. This paper asks a different and prior question. **Has anything ever actually run this problem to completion?**

One thing has, and it is not a computer system. The Theravāda Pāli canon was transmitted for roughly four centuries with no written copy at all, and for roughly two millennia in total, without a living author to arbitrate, across schism — on the Sri Lankan chronicles' account, one that produced a rival recital — through famine, war, and the ordinary attrition of the human beings who *were* the storage medium. That is a copy-integrity problem under a threat model that overlaps a retrieval server's without being the same one (§7 partitions the two), sustained for a duration no production system has approached. It converged on an apparatus, and this paper reads seven of its parts as engineering.

All seven are transpositions, and the paper says so plainly. Five transpose practices the archival and transparency-log communities already know under other names: **an inline attribution formula** opening the text rather than a metadata field (§4.1); **admission decided at build time by a named assembly** under a recorded procedure rather than by a query-time filter (§4.2); **commentary labelled as a distinct text class**, from which a server inherits an obligation never to substitute generated summary for source (§4.3); **variant readings preserved with their line of transmission attached** rather than normalised to a preferred text (§4.4); and **abridgement announced inline** at the point of the cut (§4.5). Two more transpose practices known in other fields: **interpretive standing carried per passage** rather than per document, because an excerpt loses its document's front matter at retrieval (§4.6); and **scheduled public re-verification before**

an assembly in which silence is a positive assertion for which the silent party is accountable (§4.7). An eighth item is a convention rather than a mechanism — **every exclusion served with the case that occasioned it** (§4.8).

Two further designs (§5) are carried furthest from where the destination field keeps them. The **conformance endpoint** (§5.1) accepts a claimed quotation and returns a disposition *on the claim*, together with the procedure to follow next, and by construction returns no disposition on the claimant — the part a hash does not have, because a hash reports that a comparison failed and says nothing about what to do. It is disclosed as a composition of known parts: closed disposition sets about a submitted object are old (OCSP), and what is disclosed here is a set over a multi-recension corpus that includes a conforming variant with its witness. The **separated completeness manifest** (§5.2) is a signed, versioned and expiring statement of a collection's membership and cardinality published as an artifact distinct from the index it describes, on the ground that **an index cannot attest to its own completeness**: a silently dropped document is invisible to every check that reads the index, and detectable to a third party holding only the manifest. It is a transposition, with The Update Framework's snapshot and timestamp metadata and BagIt's `Payload-Oxum` as its nearest prior art.

The paper's most important result is a narrowing, and it is stated in the abstract rather than buried in limits. **The transmission did not leave one text.** It left branches — recensions associated with different reciter lineages — and divergence between them that is detectable today because written branches and parallel traditions survived, attributable in part, and attributed largely by commentaries written centuries later. So what the two millennia left was not the prevention of divergence but divergence that can be **detected and, in part, attributed**. That is a weaker property than the one a reader might expect, and it is the *right* property for a retrieval system, whose actual failure mode is not that its copies differ but that they differ silently. It is also a retrodiction at $n = 1$ and particular to this corpus: Vedic recitation, under constraints of the same kind, aimed at and built for prevention (§9.2). §6 then argues the case against inheritance: the assembly's power to exclude, and the authority to rule on a passage's interpretive standing, are capture surfaces, and a served corpus must guard them or refuse them. §9 states the limits without hedging — survivorship at $n=1$ over two thousand years, a fitness function that plainly included royal patronage and not only textual accuracy, evidence that the apparatus itself admitted, and the observation that four of the seven mechanisms, and the eighth convention, are *social* protocols whose engineering forms are half the mechanism at most. **A scheduled re-verification with no assembly to verify before is a cron job talking to itself.** Offered under CC0 1.0 Universal as defensive prior art.

Connection to the mission frame. This institution has committed to writing its corpus primarily for a machine reader — an autonomous successor who will inherit a body of text and must be able to tell what in it is load-bearing, what is commentary, what has been superseded, and what was never checked. That commitment is empty unless the corpus can be *served* in a form that supports the distinction. The apparatus described here is the one the founder's own tradition built for keeping a text checkable across generations of carriers — a problem that overlaps this one — under conditions harder in several respects, for a reader who was also expected to carry the text forward without becoming its author. The transposition is not decoration on a technical claim. It is the reason the technical claim is the institution's problem rather than someone else's.

1 · Why this problem is the institution's problem

The corpus this paper belongs to is written, deliberately and on the record, primarily for a machine reader. That decision has a consequence which took some years to become visible: **a corpus written for a machine reader is only as good as the channel that delivers it**. A human scholar who receives a garbled quotation has recourse — a library, a colleague, a memory of having read the thing. A language model that receives a garbled quotation has none. It cannot tell a faithful copy from a plausible one, because plausibility is the only signal it has, and a well-constructed forgery is *more* plausible than an awkward original.

Provenance-Carrying Retrieval addressed the delivery layer directly and argued for an inversion: a retrieval server normally asks to be believed, and ours ships the means to disbelieve it. Text arrives bound to a content hash, an independent time anchor and a citable identifier, so the consuming agent can verify its own citation rather than trusting the channel.

That paper is complete and this one does not repeat it. What it did not ask — a question Indological and Buddhist studies approach from the other side, having long read oral transmission apparatus as a system of fidelity (§2.5) — is whether the problem has ever been *solved under load*. Every mechanism in the digital-preservation and transparency-log literatures is young. Certificate Transparency is from 2013. LOCKSS is from 1999. PREMIS is from 2005. Software Heritage began in 2016. These are excellent systems and this paper leans on several of them. But none of them has been run for two hundred years, let alone two thousand, and none has been run through a schism that forked its corpus and left both forks alive.

One protocol has. The institution's founder is a Theravāda Buddhist co-transcribing the Khmer Tipiṭaka with his father, which is how this observation was available to be made at all, and the observation is not devotional. It is that a very hard version of this exact problem was posed, was worked on continuously by large numbers of intelligent people under adversarial conditions for a duration that dwarfs the entire history of computing, and produced an apparatus that can be read off and reused.

Three things make it worth a paper rather than a remark.

First, the conditions are not a loose analogy — they overlap, and several of them are harder (§7 partitions where the two threat models do not overlap). No living author to arbitrate a disputed reading; that is the position of any corpus after its authors die, and it is the position this institution has explicitly planned for. No writing, for the first several centuries; that is a storage layer with no durability guarantee whatsoever, which is strictly worse than any modern substrate. Active schism; that is a fork with two live maintainers and no upstream, which is the governance failure mode every open corpus eventually faces.

Second, the mechanisms have engineering forms that are not metaphors. §4 states each one as a plain design rule first and gives its canonical source second, and that ordering is deliberate: the section is constructed so that **every mechanism survives the deletion of every Pāli term in it**. This is the corpus's standing deletion test applied structurally rather than as a promise — if a reader strikes the tradition out, the design rules remain, intact and independently arguable. That is a property of how the section is built, not a rule the authors ask to be trusted on.

Third, the transposition returns two designs the destination field has only in part — one a composition of known parts, one a transposition of software-update practice to a served text corpus — **and a refusal it**

does not have. A paper that only recovered known practice would be a pleasant essay. §5 and §6 are why this is filed as prior art.

2 · Background and prior art, engaged generously

The claim is transposition. That obliges an honest account of what the destination field already has, and this section is deliberately unflattering to the paper's own novelty.

2.1 — Digital preservation

PREMIS supplies a data dictionary for preservation metadata — provenance, fixity, rights — and is the standard vocabulary for saying *this object came from there and has this checksum*. **OAIS** (ISO 14721) supplies the reference model for an archive, whose concepts PREMIS builds on. **PROV-O** supplies a W3C ontology for provenance as a graph of entities, activities and agents. Between them these cover a great deal of what §4.1 describes, with one difference this paper is making the whole of its case on: **they are metadata standards, and metadata is a separate stream from the content**. An agent that retrieves a passage and does not fetch its PREMIS record has the passage and no provenance. §4.1's claim is about *where the marker lives*, not about whether provenance should exist.

2.2 — Replication, reference and identity

LOCKSS — Lots Of Copies Keep Stuff Safe — is the closest thing in the destination field to the *bhāṇaka* lineage structure, and the resemblance is real: independent institutional copies, periodic mutual comparison, repair by consensus among peers. **Memento** solves time-travel for the web, which is the *dated witness* problem. **Software Heritage** assigns intrinsic identifiers (SWHIDs) computed from content rather than assigned by an authority, which is the most philosophically similar move in the whole destination field: identity from the bytes, not from a registrar.

This paper does not claim to improve on LOCKSS. Where §4.4 differs is narrower: LOCKSS repairs toward agreement, and the mechanism in §4.4 explicitly refuses to repair, preserving disagreement with attribution instead. Those are different goals, and the canon's is the unusual one.

2.3 — Transparency logs and supply-chain attestation

RFC 6962 Certificate Transparency is load-bearing prior art for §5.2 and it must be credited hard. CT already has signed tree heads carrying the tree's size, inclusion and consistency proofs, and a gossip requirement whose mechanism RFC 6962 left to later work — which together give append-only-ness and, where gossip is deployed, detection of split views. The institution's own ledger design depends on it. **Sigstore**, **Rekor** and **SLSA** move the same primitives onto build artifacts; **SCITT** and **COSE receipts** generalise transparency receipts as an IETF work item; **C2PA** carries provenance manifests for media, including external manifests that travel separately from the asset.

The nearest prior art for §5.2 is not CT, and it is closer than CT. The **Update Framework (TUF)** separates a *snapshot* role, which signs a statement of the latest version of every targets metadata file on a repository and so prevents mix-and-match attacks, from a *timestamp* role that re-signs frequently, and it requires clients to refuse expired metadata and any version lower than one already trusted — which defeats

freeze and rollback. **BagIt** (RFC 8493) carries a `Payload-0xum`, the octet count and file count of a package's payload, "intended for the purpose of quickly detecting incomplete bags."

Against that, §5.2's claim is smaller than it might look and is stated at its true size in §5.2 itself. **CT proves that a log is append-only and that a given entry is in it; it does not prove that the log contains everything it ought to contain.** TUF and BagIt do address completeness against a declared expectation, each in its own setting, so §5.2 is offered as a transposition of that practice to a served text corpus read by machines, not as a design the destination field lacks.

2.4 — Textual scholarship

This is the field with the oldest and strongest claim, and pretending otherwise would be the paper's most embarrassing failure. **Lachmannian stemmatics** reconstructs an archetype from the pattern of shared errors among witnesses. **Greg-Bowers** copy-text theory formalises editorial choice between variant readings. The **TEI critical apparatus** is a mature XML vocabulary for encoding variants with their witnesses — which is to say, **§4.4 is a description of standard practice in textual scholarship and claims no novelty whatsoever.** *Resetting*, *witness*, *recension* and *collation* are terms of art in analytical bibliography and are borrowed here, not coined.

⚠ **The honest form of §4.4's contribution is therefore this and only this:** the destination field for retrieval-serving has not adopted the apparatus that textual scholarship has had for nearly two centuries, and machine-facing retrieval is precisely where its absence is most costly, because the consumer cannot supply the missing judgement.

2.5 — Oral tradition and the classics

Parry and Lord on oral-formulaic composition established that formulaic repetition is a *compositional* technology in oral epic. The Pāli material's heavy formulaic repetition — the reason §4.5's elision marker exists at all — sits squarely in that lineage.

Reading an oral canon's apparatus as a system of fidelity is not new either, and this paper's first claim is narrowed accordingly. Staal (1986) read Vedic recitation — with its forms word by word, in overlapping pairs and in dense interlaced sequences — as a fidelity of oral tradition; Cousins (1983) examined the Pāli material as oral literature, and Wynne (2004) its oral transmission; Collins (1990) read the idea of a closed Pāli canon historically, as the product of a particular monastic lineage; von Hinüber (1996) surveyed Pāli literature as a whole. The Buddhist-studies literature on the councils, the *bhāṇaka* lineages, and the relation between the Pāli recensions and the surviving Sanskrit and Chinese Āgama parallels is very large and this paper does not contribute to it. **What this paper adds is the transposition only: the mechanism-by-mechanism move from that apparatus to machine-facing retrieval serving.**

2.6 — Where the transposition adds nothing

Stated plainly, because a paper that finds itself novel everywhere is not to be trusted:

- **§4.1 (inline marker)** adds a *placement* argument to a solved provenance problem. If your consuming agents reliably read metadata, this mechanism is worth nothing to you.

- §4.2 (**build-time gate**) is admission control. Every curated database has one. The contribution is the argument for *when* it runs, not that it exists.
- §4.3 (**commentary as a class**) is old practice under a class label — the printed Talmud page, the Tibetan canon's division into Kangyur and Tengyur — and its non-summarisation obligation is already argued in *Provenance-Carrying Retrieval*. No novelty.
- §4.4 (**variants with witnesses**) is TEI. No novelty at all.
- §4.6 (**per-passage standing**) is anticipated at that granularity — per statement in RFC 2119, per report in hadith grading, per point of law in legal citators. What remains is only its symmetric default.
- §4.8 (**exclusions with reasons**) is close to standard practice in well-run archives and is identical in spirit to a good deprecation policy.
- §4.7 (**scheduled re-verification**) is a health check with an audience. Whether the audience is a real addition is argued in §4.7 and is the least secure claim in §4.
- §5.2 (**completeness manifest**) is The Update Framework's snapshot metadata and BagIt's `Payload-Oxum`, carried to a served text corpus.

What is left after that subtraction is §4.6's symmetric default, the composition in §5.1, the refusal in §6, and the transmission finding in §3.3 — with §5.2 as the transposition most worth making. That is the paper.

3 · What the field test actually was, and what it left

3.1 — The conditions

Stated as a systems problem, without devotional framing:

THE TRANSMISSION PROBLEM, AS POSED

authoritative replica	none after the author's death
storage medium	human memory (c. 4 centuries), then palm leaf, then print, then bits
durability of a node	one human lifetime
write path	communal recitation, in assembly
partition events	war, famine, monastic schism
adversary	time, politics, honest misremembering, and schism (on the chronicles' account, a rival recital with its own carriers)
duration	~2,000 years, continuous

COMPARE: the longest-running production transparency log
is younger than most of the people reading this.

The reciter lineages — *bhāṇaka* [IV] — divided the collections between them, one lineage specialising in the long discourses, another in the middle-length, others in the numerical and connected collections. This is documented independently of the tradition's own account: Sri Lankan cave inscriptions dated between the third century BCE and the first century CE name monks by the collection they carried [IV].

Two properties of that arrangement matter to a systems reader. It is a **shard**, so no single carrier held the whole corpus and no single death lost it. And it is a **replication group per shard**, since a collection was carried by a lineage rather than an individual, with recitation in company as the comparison operation. Death was not merely attrition; it was **the event that forced the verification to run**. A text living in bodies is re-verified on a schedule set by funerals, and it is self-correcting *because* it is fragile. That observation belongs to the essay *The Last Carrier* and is cited here rather than re-argued.

3.2 — The write path was consensus, and it ran at build time

The councils — *saṅgīti* [V], literally *reciting together* — are the corpus's build events. The procedure is recorded, and it is more specific in the secondary accounts than in the canonical one, a difference this paper states rather than smooths. **The canonical account of the first council** (Cullavagga XI) records a question-and-answer recitation before the assembly: under a formal motion, the convening elder asks one named monk about the Vinaya — for each rule, where it was laid down, its origin story, the person, the rule, its additions, the offence and the non-offences — and another named monk about the discourses; and the assembly's formal decisions in that account are carried by motion, with silence as assent and any monk who does not approve required to speak [V]. **The secondary accounts add that an elder recited a passage, the assembly chanted it back in chorus, and material was compiled only on unanimous acceptance.** That characterisation is theirs; it was not found in the canonical passage this paper checked (§11).

⚠ **The historicity of the first council is contested in modern scholarship, and the traditional dating is not accepted outside the tradition.** This paper does not need it to be historical. What it needs — and what is not seriously disputed — is that the *procedure* is the one the tradition recorded as normative and subsequently followed, because the engineering claim in §4.2 is about the shape of the procedure, not about the events of any particular century.

3.3 — ★ What it left, and this bounds everything else

Here is the result that a reader should carry away even if they take nothing else.

The transmission did not leave one text. It left branches. Different *bhāṇaka* lineages appear to have exercised independent judgement over the arrangement of their collections and over which versions of stories and doctrines they preserved, and variant readings between the long and middle-length collections may be attributable to different reciter schools [V]. Those lineages later developed into distinct interpretive schools whose differing views are visible in the commentarial literature [V]. On the Sri Lankan chronicles' account, the losers of the second council held a rival Great Recital of their own [V]; modern scholarship separates that council from the later schism.

A naïve reading of this paper's thesis would take that as refutation. It is not, and the reason is the paper's central finding:

What two millennia left is divergence that can be detected and, in part, attributed — not one undiverged text, and not the prevention of divergence.

⚠ **That is what the transmission left, not what it was built to do, and the paper does not present it as a selection result.** Every mechanism described here aimed at fidelity; divergence survived despite them.

What makes the divergence detectable today is the survival of written branches, of parallel traditions — the Sanskrit and Chinese Āgama parallels — and of modern collation; and the attribution of readings to lineages is made largely by commentaries written centuries after the fact. **The finding is a retrodiction at n = 1 (§9.2).** Its engineering conclusion does not depend on reading it as selection.

The variants survive, and much of the time their line of transmission can still be read. The disagreement is legible. A modern editor can collate the branches and say where they differ and, often, whose transmission a given reading came through. That is a completely different outcome from a corpus that silently converged on one reading and lost the record of what it overwrote.

And it is the property a retrieval system actually needs. **A server's failure mode is not that its copies differ. It is that they differ silently** — that a normalisation, a re-encoding, a well-meaning correction or an outright substitution passes through the layer and arrives looking exactly like the original. A guarantee of correctness is not on offer from any system, including this one. A guarantee that divergence is *visible and attributable* is on offer, and it is what these mechanisms deliver.

🚫 **The consequence binds the rest of the paper and is repeated in §9.1: every integrity claim here is a detection claim. None of these mechanisms prevents a corrupted text from being served, and the paper must not be cited as though it did; the design rules say what a server must never do silently.**

4 · The mechanisms, stated as design rules first

★ Every subsection below states the engineering rule before naming its canonical source, and the rule is written to stand without the source. Strike every Pāli term from this section and the design survives entire. That is the corpus's deletion test satisfied by the section's construction rather than by the authors' assurance — a property, not a promise.

MECHANISM	ENGINEERING FORM	NOVEL?
1 opening attribution formula →	provenance inline in the text stream, not in metadata	no (placement arg.)
2 assembly recitation →	admission decided at BUILD time by a named body	no (timing arg.)
3 commentary as a text class →	never return summary where source is expected	no (class label: old)
4 recensions kept distinct →	variants served with their witness; never normalise	no – this is TEI
5 announced abridgement →	elision marked inline, at the cut, and recoverable	no
6 interpretive standing →	standing per PASSAGE, since an excerpt loses front matter	partly (symmetry only)
7 scheduled re-verification →	public, before an assembly, silence load-bearing	weakest claim in the set
+ recorded exclusion reasons →	every exclusion published with its occasioning case	no
AND SEPARATELY, §5:		
the conformance endpoint	– a composition	
the completeness manifest	– a transposition (TUF; BagIt)	

4.1 — Provenance travels *in* the text

The rule. A unit of served text carries its attribution as part of the text itself, at its head, so that any excerpt of it that survives a copy-paste, a context window, a summarisation or a quotation still carries the attribution. Provenance recorded only in an accompanying field is provenance that a consumer may decline to fetch, and machine consumers decline constantly — not maliciously, but because the field was not in the retrieved span.

The canonical instance. The discourses of the principal collections open with an attribution formula, conventionally rendered *thus have I heard*, followed by the setting and the occasion [V] ; other collections carry their own — the Itivuttaka's discourses open *this was said by the Blessed One* [V] — and in every case the formula is inline, never a header. It is the first sentence of the text, in the same voice, and it is preserved through every recitation and every copy because removing it would be removing text.

Against §2.1. PREMIS, PROV-O and OAIS all solve provenance *representation* better than this does. The narrow claim is placement, and the argument for it is the one in §3.3: a metadata record that the consumer did not fetch does not fail loudly. It fails by absence, and absence is the failure mode this whole paper is about.

Nearest prior instances. The hadith *isnād* is the strongest: a chain of transmitters carried in the same stream as the report it vouches for, placed before the text, and the basis on which the report is graded. In retrieval practice, contextual chunk-prefixing (published by Anthropic as *Introducing Contextual Retrieval*, 19 September 2024) prepends a statement of document context to each chunk before it is indexed — inline, as this rule asks, but generated rather than attributed, which is exactly the difference the rule turns on.

Sig-9 check — property or rule? Inline is a **property**: the marker survives excerpting because of where it sits, with no enforcer. A metadata field is a **rule**: it survives because something remembered to fetch it. This is the clearest instance in the paper of the institution's own take-the-fix-from-the-object test, and it is why §4.1 leads.

4.2 — Admission is decided at build time, by a named body, under a recorded procedure

The rule. What is in the corpus is settled when the corpus is assembled, by an identified group following a procedure that is written down, and the decision is recorded. It is not settled at query time by a filter. A query-time gate has an operator, and an operator can be leaned on per request; a build-time gate produces an artifact that either contains a document or does not, and the containing is checkable by anyone.

The canonical instance. The councils, whose procedure §3.2 describes: in the canonical account, recitation by named monks in answer to the convening elder, before an assembly whose formal decisions are carried by motion with silence as assent [V] ; in the secondary accounts, a chorus and compilation **only on unanimous acceptance**. The rule needs only what both accounts share — a named body, a recorded procedure, a decision made at assembly.

What this is not. It is not a claim that assemblies produce correct results. §6.1 argues the opposite risk at length. It is a claim about *where the decision sits* and therefore about what can be audited afterward.

Relation to our own corpus. *Buddha AI and the Living Tipiṭaka* §8 addresses the adjacent question of who may canonise *new* material and specifies a tiered workflow and periodic councils for it. That is a different question from this one: §8 asks what may enter a canon, this asks how a *served copy* of an existing one is assembled. Both are cited; neither absorbs the other.

4.3 — Commentary is a distinct class, and the server therefore never summarises

The rule. If a corpus distinguishes primary text from commentary on it, a retrieval layer over that corpus must serve them as distinct classes and must never return generated prose in the position where primary text is expected. A summary is not a shorter version of a text. It is a *different text*, of unknown provenance, and it cannot be hash-verified against anything.

The canonical instance. The commentarial layer — *aṭṭhakathā* — is labelled as a separate class of literature with its own transmission and its own standing [V]. ⚠️ **The label held; the boundary did not.** The canon itself contains texts of commentarial form: the Niddesa, a commentary on parts of the Sutta Nipāta, sits in the Khuddaka Nikāya [V], and the Vinaya's old word-commentary is incorporated into the Suttavibhaṅga [V]. The Burmese Khuddaka admits the Nettipakaraṇa, the Peṭakopadesa and the Milindapañha, where the Sinhalese printed edition admits the first two [V]. And the earlier Sinhala commentaries from which the Pāli *aṭṭhakathā* were compiled are lost [V] — the redaction survived and its sources did not. What the tradition kept is an explicit **class label**, maintained even while material crossed the line it drew.

The class label is old practice elsewhere. The standard printed Talmud page (the Vilna edition's layout) sets the Mishnah and Gemara in the centre column with Rashi's commentary and the Tosafot in columns around it, visibly distinct; the Tibetan canon divides the translated word of the Buddha (Kangyur) from the translated treatises and commentaries (Tengyur).

Novelty: none. The non-summarisation obligation is stated in *Provenance-Carrying Retrieval* as a consequence of hashability, and the engineering rule stands on that argument alone: a summary is a different text of unknown provenance and cannot be verified against anything. The canonical instance adds no independent reason for it. What it adds is a caution: a class boundary moves over time — texts are admitted, absorbed and reclassified — so a server should serve each text's class label together with its reclassification history under §4.8, as §6.2 asks for the standing field, rather than treat the boundary as fixed.

4.4 — Variants are served with their witness, never normalised

The rule. Where the corpus has divergent readings, serve them as divergent, each with the line of transmission it came through. Do not silently pick one. Do not merge them into a smoothed text. A consumer that receives a normalised reading has been told a lie of omission that it has no way to detect.

The canonical instance. The recensions associated with different reciter lineages, whose variant readings may be attributable to those lineages [V], and whose differences the commentarial literature records rather than resolves [V].

⚠️ **Novelty: none.** This is the TEI critical apparatus, and before it Lachmann and Greg–Bowers. §2.4 says so. The only contribution is the observation that machine-facing retrieval has not adopted it and is the

setting where its absence costs most — because a human reader of a smoothed text can suspect smoothing, and a model cannot.

4.5 — Abridgement is announced inline, and is recoverable

The rule. When served text is abridged, mark the cut *at the cut*, in the text, in a form that names the elided span by a resolvable identifier or digest sufficient to fetch and verify it. Do not record the abridgement in a field. Do not abridge silently, ever.

The canonical instance. The elision marker — *peyyāla*, abbreviated in practice to a short particle [V] — which stands in the text where a formulaic passage recurs, indicating that the passage is to be supplied. ★

The precision matters and this paper states it rather than glossing over it: the marker does not signal *arbitrary omission*. It signals a *repetition to be supplied* — from the surrounding text, where enough of the pattern is given, or, where it is not, from a formula the reciter already held or from another text [V] ; a modern translator of the canon calls these *internal* and *external* abbreviations. A sample of every elision in one long discourse (the Mahāparinibbāna Sutta, twenty-nine instances in the root text) found each abbreviating repeated or formulaic material and none marking an arbitrary cut. An elision marker that says only "something was here" is decoration; one that says "the passage you already have, repeated" is a compression scheme with an integrity property.

⚠ **The external case is the one a retrieval system inherits badly.** The reciter already held the formula. A reader who has fetched one passage from a served corpus holds nothing, which is why the engineering form cannot stop at a marker and must name what was cut resolvably.

The engineering form is therefore narrower than "announce elisions": announce them in-band, at the point of the cut, with an identifier or digest from which the elided span can be fetched and verified. That is a real constraint and most truncation in retrieval systems fails it.

Nearest prior instance. TEI's <gap> element marks an omission in a transcription in-band, with attributes for the reason and the extent of what was omitted; its standard attributes record why and how much, not where the omitted text can be fetched.

4.6 — Interpretive standing is carried per passage

The rule. Whether a passage should be read as stating its meaning directly, or as requiring inference, is a property of the passage. It must travel with the passage. Document-level front matter does not survive retrieval: a system that returns an excerpt has, by construction, dropped the document's `status`, `revised` and `genre` fields, and the excerpt now reads as though it carried the document's authority.

★ **The rule is symmetric, and for an engineering reason.** The instinct is to worry only about over-claiming — treating a hedged passage as settled. But a standing field that marks every passage *requires inference* carries no information, and it misstates every passage that is in fact settled: it is a misclassification in the other direction, not a safe setting. A retrieval layer therefore may not mark everything provisional as a defensive posture. **Blanket hedging is a correctness failure, not a safe default.**

The canonical instance. The distinction between a discourse whose meaning is explicit — *nītattha* — and one whose meaning is in need of interpretation — *neyyattha* [V] . ⚠ **The canonical unit is the discourse**

(*suttanta*), **not the passage**. AN 2.24 names the two who misrepresent the teacher — one who explains a discourse in need of interpretation as explicit, and one who explains an explicit discourse as in need of interpretation — and AN 2.25 names their correct counterparts [V]. The canon thus treats misclassification in either direction as one error, which is an instance of the symmetry above; **the per-passage granularity is this paper's design, not the canon's**, and its reason is the engineering one already given: an excerpt loses its document's front matter.

Nearest prior instances. Standing below the level of the document is established practice in other fields: RFC 2119 fixes a requirement level per statement (*MUST*, *SHOULD*, *MAY*) inside a specification's own text; hadith scholarship grades each report (*ṣaḥīḥ*, *ḥasan*, *ḍaʿīf*) rather than each collection; and legal citators attach treatment to a cited authority — KeyCite down to the headnote, the individual point of law. What §4.6 adds is only the argument from excerpt loss for a served corpus and the symmetric default.

Locus note. The pair is cited here in SuttaCentral's numbering, AN 2.24–25 (within its range AN 2.21–31); enumeration of this collection's early books differs between editions and a reader may encounter it under another number. See §11.

4.7 — Scheduled public re-verification, in which silence is load-bearing

The rule. Integrity is re-affirmed on a fixed, public, predictable schedule, before an assembled audience rather than to a log. The affirmation is solicited a set number of times. **Silence constitutes a positive assertion of integrity, and the silent party is accountable for it.**

The canonical instance. The fortnightly recitation of the disciplinary code — *pāṭimokkha* — before the assembled community on the full and new moon [V]. The community is told that one who has committed an offence should declare it and one who has not should remain silent, and that *from your silence I shall understand that you are pure* [V]. The announcement is made up to a third time, and **one who remembers an offence and does not reveal it is lying in full awareness** (*sampajānamusāvāda*), which the accompanying word-commentary classifies as an act of wrong conduct, a *dukkata* (Mahāvagga II.3.3, II.3.7) [V].

★ ★ **That last clause is what makes the mechanism more than a health check, and it is the reason this subsection exists.** An unanswered health check is ambiguous: the service may be fine, or unreachable, or lying. Here silence is *defined* as an assertion and is *penalised* if false, which converts an absence into a signal with a cost attached. A monitoring system that treated "no response" as "healthy" would be badly designed; a monitoring system in which "no response" is a *signed* claim of health is a different thing entirely.

⚠ **This is the weakest claim in §4 and the paper marks it as such.** The engineering form — periodic signed attestations of integrity, published, with liability for a false attestation — exists in the transparency-log world already, and the addition of *an assembly* is precisely the part §9.3 argues may not transpose at all.

Nearest prior instances, and they are older than the transparency log. Silence as assent is ordinary practice: an auditor's *negative* confirmation request is answered only by a party who disagrees, and a non-response is treated as agreement; a chair who asks *is there any objection?* takes silence as unanimous consent. A signed periodic attestation with a cost for a false one is ordinary too: under the Sarbanes-Oxley

Act, named officers certify each periodic report (§302), and a knowingly false certification carries criminal penalties (§906). §4.7 combines the two; neither half is new.

4.8 — Every exclusion is served with its recorded reason

The rule. A corpus that excludes material publishes, alongside each exclusion, the case that occasioned it — and the subsequent amendments and exceptions, so the exclusion set has a history rather than a state. An exclusion list is an assertion; an exclusion list with reasons is auditable.

The canonical instance. Each disciplinary rule is preceded by its origin story — the specific incident that prompted it — and followed by the amendments the rule accumulated as new situations arose, together with exceptions marking what is *not* a violation [V] . Further cases function as judicial precedent [V] .

★ **The amendment history is the part the destination field usually drops.** Archives commonly publish a takedown policy; they rarely publish, per removed item, the case that caused the rule to exist and the successive refinements to it. That is the difference between a policy and a case law, and a corpus that will outlive its curators wants the latter.

Nearest prior instances. The Lumen database publishes copies of content-removal requests that online services have received; a rejected CVE record stays on the CVE List marked *REJECTED*, most often with its reason stated in the record's description. Both publish an exclusion together with its occasion. Neither is organised as a case law, in which a rule's later amendments and exceptions accumulate against the case that first prompted it.

5 • The two designs carried furthest from the destination field


§2.6 subtracted the paper down to what is left. Both designs below take their shape from a canonical mechanism, and what is claimed is their placement in a retrieval server, not their invention — the first as a composition of known parts, the second as a transposition with close prior art, each named where it is stated.

5.1 — ★ The conformance endpoint

The problem a hash does not solve. A machine reader that verifies a quotation against a content hash learns exactly one bit: matched, or did not. That bit is nearly useless in the case that actually occurs. A model quotes a passage it half-remembers, or a user pastes a quotation from a third-party site, or an agent carries a citation forward from a prior turn. The hash fails. And now what? The failure is indistinguishable between *the quotation is fabricated*, *the quotation is real but from a different recension*, *the quotation is real and this corpus has been tampered with*, and *the quotation is a faithful paraphrase that was never a quotation at all*. **A hash reports that a comparison failed. It says nothing about what to do next, and it therefore cannot prevent the single most common downstream behaviour, which is to shrug and use the text anyway.**

The mechanism. An interface that accepts a *claimed* text — not a hash of it, the text — lays it alongside the served corpus, and returns a **disposition on the claim** together with the procedure the requester should

follow. The disposition set is small and closed, and the point of the design is that each member names a different next action:



Mapped to the four failures a hash cannot tell apart: a fabricated quotation returns *NOT FOUND*; a real quotation from another recension returns *CONFORMS, VARIANT* with the witness it came through; a faithful paraphrase returns *NOT FOUND* — it is not a quotation from this corpus and is not to be quoted as one; and tampering with the served corpus itself is not detectable by the endpoint at all, since the endpoint compares against that corpus, and is caught instead by the corpus's own hash and time anchor (*Provenance-Carrying Retrieval*).

The canonical source. The four great references — *mahāpadesa*, given in the account of the teacher's last days (DN 16) and again in the numerical collection (AN 4.180) [V] — instruct that when someone claims to have received a teaching, the statement is to be neither approved nor rejected, but its words **carefully learned and checked against the discourses and the discipline** [V]. If they do not fit, the conclusion is stated in two parts: this is not the teacher's word, and *it has been incorrectly memorised* — by that monk, that community, those elders or that elder, according to the source the claimant named — and it is to be discarded [V]. ⚠️ **So the canonical disposition does name the carrier: as the one whose memorisation erred, with no penalty prescribed.** It is a finding about the fidelity of a copy rather than a verdict on a person — the canon already blames the copy, not the teaching — and the test it prescribes is the *accordance of a teaching* with the collections, not the collation of a string. **Two steps here are this paper's, not the canon's:** the transposition from doctrinal accordance to quotation collation, and the removal of the carrier from the disposition altogether.

☆☆ **That second step is the design.** A verification interface that returns a verdict on the *claimant* creates an incentive not to submit claims, which destroys the mechanism's coverage exactly where coverage matters. A verification interface that returns a verdict only on the *claim* is safe to use, and therefore gets used. It is a small distinction and it is the entire difference between a conformance service and a reputation service.

What is disclosed, and what is prior art. A closed disposition set about a submitted object is not new: OSCP (RFC 6960) returns *good*, *revoked* or *unknown* about a submitted certificate and says nothing about the party asking. Matching a claimed reference against a record is not new either: Crossref's bibliographic query takes a free-text citation and returns the works it may refer to. Hash verification is universal, and so is text similarity search. C2PA's validation reports say what could and could not be verified, but of an asset's

own manifest rather than of a third party's claim about a corpus. **What is offered is the composition:** a closed disposition set over a multi-recension corpus that includes *conforms, as a variant, with its witness*, each member paired with a prescribed next action, and no place in the response schema for the claimant.

Sig-9 check. The claim/claimant separation is a **property** if the disposition set has no member that can carry a verdict on a person, and a **rule** if the operator is merely asked not to add one. It must be built the first way. 🚫 *Remove the enforcer: if the schema permits a claimant field, someone will populate it under pressure, and the mechanism becomes a denunciation channel.*

⚠️ **The schema is not the only surface, and the same test applies to the log.** An endpoint that accepts text over a network receives every request with an address, a key or an account. If it retains the requester's identity with the submitted text, the record of who submitted what — the denunciation record — exists on the server whatever the response schema says, and the no-claimant property has become a rule about the log. **It is a property only if the endpoint retains no requester identity with the submitted text.** And even then the schema stops the *server* from judging a claimant; it cannot stop a consumer from republishing a disposition — *NOT FOUND*, about a quotation a named person circulated — as a judgement of that person.

5.2 — ★ The separated completeness manifest

The problem transparency logs do not solve. Certificate Transparency and its descendants prove that a log is append-only and that a given entry is in it. **They do not prove that the log contains everything it ought to contain.** An operator who never adds a document has added no dishonest entry; every inclusion proof still verifies; every consistency proof still verifies. Omission is invisible to the entire apparatus, because the apparatus reads what was *produced*, and a document that was never produced produces nothing to read.

⚠️ **This is the sharpest form of an instrument failure the institution has recorded repeatedly in its own operations: every check reads an artifact, so no check can report a thing that was never made.** An unstamped document, a workflow that never ran, a field nobody wrote — each is invisible to the check that would have caught it. The log is the same shape.

The mechanism. Publish, as an artifact **distinct from the index**, a signed enumeration of a collection's membership and its **cardinality**, so that a third party holding the manifest alone can detect a document's silent removal. The separation is the mechanism, not an implementation detail: **an index cannot attest to its own completeness**, for the same reason a witness cannot corroborate themselves. If the manifest lives inside the index, an operator who removes a document removes its manifest line in the same edit and nothing anywhere disagrees.

The manifest carries a monotonically increasing version and an expiry, and a holder refuses an expired manifest or one older than a version it has already seen. Without both, a stale manifest is internally consistent: an operator who reverts the collection to an older, smaller state and serves the older manifest that matches it passes every check that reads the pair (§10, *the frozen manifest*).

⚠️ **Out of band here, in band in §4.1 — and the two are not in tension.** §4.1 puts provenance in the text because its consumer is a reader who never elected to fetch a separate record, and whose failure when the

record is missing is silent. The manifest is kept out of band for *independence*. Its holder is an auditor who elected to hold it, and for that auditor its absence is not silent.

The canonical source. The mnemonic summary verses — *uddāna* [V] — placed at the end of each section and of the whole work. Their function is precisely this: they key off each member of a group, fixing **membership and ordering** [V], and the summaries at the end of a work give **the names and the counts** — of the elders, of the verses in each chapter, and of the whole [V]. They were memorised alongside the text by the reciters [V].

⚠ **The best-known instance does not balance, and §9.11 treats this as a limit on the paper.** The totals that the Theragāthā's summaries give do not match a plain count of its verses, and the plain count itself differs between editions.

★ **Two details of the canonical form are load-bearing and both were confirmed rather than assumed.** The summaries carry **counts**, not merely names — cardinality is what makes a removal detectable when a name is also removed. And they are **separate units**, positioned at the boundary of the collection rather than distributed through it, which is what allows a carrier to hold the manifest without holding the collection.

Against §2.3 — and it is a transposition. CT gives append-only-ness and split-view detection; a completeness manifest gives completeness against a *declared* expectation. They compose, and neither substitutes for the other. But the manifest is not a new design — §2.3 names its nearest prior art in The Update Framework's snapshot metadata and BagIt's `Payload-0xum` — and **this paper withdraws any claim that the design is absent from the destination field**. What it offers is the transposition to a served text corpus read by machines, together with the canonical instance of keeping the count apart from what it counts. The honest statement of the residual is that the manifest moves the trust problem rather than dissolving it — a signed manifest is only as good as the key that signed it and the party that published it — and §9.5 says so.

Sig-9 check. Separation is a **property** when the manifest is signed by a different key and published by a different party than the index. It degrades to a **rule** the moment the same operator holds both, at which point it detects accident and not intent. ⚠ *This is a real limit and it means the manifest is worth much less inside a single-operator system than it looks.*

6 · What must not be inherited

🚫 **A transposition that finds every mechanism inheritable is doing apologetics.** Two of the §4 mechanisms carry capture surfaces, and naming them is claimed in §Claims 13 as a design element of this paper rather than as a caveat on it.

6.1 — The assembly's power to exclude

§4.2 argued for a build-time gate operated by a named body. The same mechanism, viewed from one step back, is **an assembly with the power to decide what does not exist**. On the Sri Lankan chronicles' account, the losers of the second council held a rival Great Recital of their own [V] — modern scholarship separates the council from the later schism — and either way a schism is what an exclusion looks like from the excluded side.

This institution's own doctrine refuses exactly this shape elsewhere: a guard that depends on the good behaviour of whoever holds it at the moment it is tested is a rule, and rules need a living enforcer with the right incentives at the exact moment nobody can guarantee. **An admission body is a rule wearing a procedure's clothes.**

★ **The available guard is §4.8, and it is only partial.** Serving every exclusion with its occasioning case and its amendment history makes the exclusion set auditable, which raises the cost of a bad exclusion without preventing one. **We state plainly that we do not have a property-shaped fix for this and are not claiming one.** The honest position is that a served corpus with a curated boundary has a governance problem that no amount of cryptography touches, and that the corpus should say so on its own surface rather than let the presence of hashes imply otherwise.

6.2 — The authority to rule on interpretive standing

§4.6's per-passage standing field is useful and is also a lever. Whoever sets the field decides which passages are to be read as they stand and which are open to inference — and the second category is where interpretation lives. An operator who can mark a passage as requiring inference can, without altering a byte, change what the corpus is understood to say.

⚠ **Two guards, and the paper is honest that they are mitigations rather than solutions.** The field must be **inside the hashed text**, so that changing it drifts the hash and rotates the proof, making the change as visible as an edit — because it is an edit. And the corpus must serve the standing field's own **change history** under §4.8, so a reclassification is as auditable as a removal.


Related and separate: the guard that a reciter is not the author of what is recited is claimed and argued in *The Song That Is Not His* and is **cited here rather than re-claimed**, per the placement rule. It bears on §4.6 in an obvious way and this paper adds nothing to it.

6.3 — The general form

Both refusals have the same shape and it is worth stating once: **the mechanisms that handle bytes transpose cleanly, and the mechanisms that handle authority do not.** Inline markers, manifests, variant witness *records*, announced elisions and conformance dispositions are all properties of artifacts. Witness *admission* — deciding what counts as an attested line of transmission — is not, and returns to §6.1, as §10 says of variant proliferation. Admission and interpretive standing are properties of a body politic, and importing them into a served corpus imports a governance question that the corpus's cryptographic apparatus is entirely powerless to answer — while, dangerously, *looking* as though it has answered it.

7 • The threat-model partition

The transposition's most likely misuse is to treat these mechanisms as a complete integrity story. They are not, and the reason is a clean partition that the paper states early enough to be useful.



An enumeration detects omission and cannot detect tampering. A hash detects tampering and cannot detect omission. They are complements, and a system with only one of them has a hole shaped exactly like the other. That partition is the practical reason the two designs of §5 are worth having *together with* the envelope of *Provenance-Carrying Retrieval* rather than instead of it.

⚠️ The partition also bounds the field-test claim. The canon's mechanisms were developed against *its* threat model, not ours. Where the two models overlap the transposition is well-motivated; where they do not, the canonical practice is evidence of nothing about our case, and §4 marks which is which.

8 • Where this sits against the rest of this corpus

Stated explicitly because a corpus that cross-references loosely is a corpus whose reader cannot tell what is claimed where.

Paper	Its question	Why this one is not it
<i>Provenance-Carrying Retrieval</i>	how does a response carry the means of its own falsification?	that paper builds the envelope ; this one asks what a long field test left, and supplies two mechanisms the envelope does not contain (§5)
<i>Buddha AI and the Living Tipiṭaka</i> §8	what makes a canon authoritative ?	authority over <i>new</i> material, decided by a doctrinal body; this paper is about the integrity of a copy of existing material
<i>The Borrowable Standard</i>	what does a corpus need before it can exist?	orthography and encoding — upstream of this paper and disjoint from it
<i>The Song That Is Not His</i>	who is the author of a recitation?	owns the reciter-is-not-author guard; cited in §6.2, not re-claimed
<i>The Last Carrier</i> (essay)	what has carried this canon, and what carries it now?	the founder's voice, the media chain, and the observation about fragility that §3.1 cites
<i>The Persistence Architecture</i>	can an apparatus be built so its author can stop?	succession, not transmission integrity

★ The residue owed elsewhere. *Provenance-Carrying Retrieval*'s prior-art section currently cites PREMIS, PROV-O, LOCKSS, Memento and Software Heritage and stops, on the line that the archival community has

thought about this longer than the machine-learning community has. **The extension — that an oral canon thought about it longer still — is a two-sentence cross-reference belonging to that paper, not a section transplanted from this one.** It is queued against that paper's next revision.

9 • Honest limits

This section makes no appeal to the tradition's authority; historical facts about the transmission appear in it only as evidence against the paper. That is deliberate and is the corpus's standing rule for honest-limits sections.

9.1 — Every integrity claim here is a detection claim

Restated from §3.3 because it is the limit most likely to be lost in citation. **None of these mechanisms prevents a corrupted text from being served; the design rules say only what a server must never do silently.** They make divergence detectable and, where a witness exists, attributable. A corpus running all of them can still serve a wrong text; what it cannot easily do is serve a wrong text *silently*. Any use of this paper that implies a correctness guarantee is a misuse, and the paper's own field test is the evidence — the tradition ran this apparatus for two millennia and ended with branches, not with one text.

9.2 — Survivorship, at $n = 1$, over two thousand years

We observe the canon that survived. The claim that these mechanisms *caused* its survival is not testable: there is no control and no record of the traditions whose apparatus failed. **This is unfalsifiable at $n = 1$ and the paper does not argue around it.**

There is, however, a comparable case, and it cuts against the paper. Vedic recitation, an older oral-conservation system under constraints of the same kind, aimed at *prevention* and was built for it: redundant recitation forms — word by word (*padapāṭha*), in overlapping pairs (*krama*), in dense interlaced sequences (*ghana*) — whose purpose is that not a syllable be altered, read by Staal (1986) as a fidelity of oral tradition. Hadith transmission, for its part, built a developed oral provenance apparatus in the *isnād*. **So detection-not-prevention is what *this* corpus shows, not what §3.1's constraint box predicts,** and Claim 2 is bounded by that independently of $n = 1$. And the transmission finding of §3.3 is itself a retrodiction at $n = 1$ — read off the branches that survived, with the attribution of readings supplied largely by later commentaries — not a selection result.

What survives both objections is narrower and is what §Claims 1 actually says: the mechanisms are *legible as engineering*, transposable, and have forms worth using on their own merits. Their canonical provenance is a reason to look, never evidence that they work.

9.3 — The fitness function was not textual accuracy alone

Sharper than survivorship, and it is the objection we would raise against this paper. The councils had royal patronage; the collections that survived are substantially the collections that had institutional and political backing. **A field test whose fitness function included state support may have selected for political durability, with the textual mechanisms partly along for the ride.** We cannot separate the two. The consequence is stated as a bound on the whole paper: what is offered is **provenance of design — an**

account of where these ideas come from and why they are worth considering — and never proof of efficacy.

9.4 — Four of the seven, and the convention, are social protocols, and their engineering forms are half the mechanism

§4.2, the witness admission of §4.4, §4.6, §4.7 and the §4.8 convention all presuppose a body of people: an assembly to admit, a body to decide which lines of transmission count as witnesses, an authority to assign standing, a community to recite before, a tradition to record precedent. **A scheduled re-verification with no assembly to verify before is a cron job talking to itself** — and it is worth saying that this is the paper's own strongest objection to itself, not one an outside reader had to supply.

This bites hardest on §4.7, which is why that subsection is marked as the weakest claim in §4. It bites least on §4.1, §4.4's witness record and §5.2, which are properties of artifacts and need no one present. **The ranking is the useful output: prefer the mechanisms that survive the removal of the community, and treat the rest as governance proposals wearing engineering clothes.** For this institution the assembly in question is a designed one whose seating is argued elsewhere and is not assumed here.

9.5 — The manifest moves the trust problem rather than dissolving it

§5.2's separation is only as good as the independence of the two publishers. In a single-operator deployment — which is what `corpus.333.eco` currently is — the same party holds the index and signs the manifest, so the mechanism detects **accident and drift, not intent**. That is worth having and is far less than it sounds. The design reaches its full strength only with an independent signer, and we do not have one.

9.6 — What is actually deployed

Stated so the reader can weigh the rest. `corpus.333.eco` is live and serves the institution's corpus with a provenance envelope. **Of the mechanisms in this paper, the ones with running code are §4.1's inline markers and parts of §4.3's class separation.** §4.2 runs as a build step with a single human operator, which satisfies its letter and not its spirit. **§5.1 and §5.2 — the two designs of §5 — are not built.** They are specified here at claim scope for exactly that reason. The Tipiṭaka server that would exercise the full set has no date.

🚫 **A reader should treat every efficacy statement in this paper as design reasoning, not as a report of operation.**

9.7 — The citation standard, stated rather than assumed

Every canonical locus is marked [V] where it has been checked against multiple independent public reference works and translations. **No locus has been checked against the Pāli by a reader of Pāli**, and the authors do not read Pāli. §11 gives the full status and names the check that is owed. A paper about checkable provenance that left its own sourcing standard implicit would be refuted by its own thesis.

9.8 — The coherence is not evidence

The mechanisms fit together well. Eight parts of an old apparatus mapping cleanly onto a modern problem is exactly the kind of pattern that is either deeply right or deeply seductive, and the two are

indistinguishable from inside. **Each mechanism is individually arguable and should be argued individually; the elegance of the set earns an experiment and does not replace one.** 🛑 The fit between the parts is not evidence for any of them.

9.9 — The evidence is self-attested

Everything this paper knows about the apparatus comes from documents the apparatus itself admitted. The first council's procedure is known from Cullavagga XI, a text transmitted by that council's successors; the attribution of readings to reciter lineages comes largely from the commentaries; the summary verses of §5.2 are known from inside the collections they summarise. **That is §5.2's failure mode applied to the paper's own evidence: a witness cannot corroborate itself**, and an account of a completeness mechanism preserved by the collection it kept inherits the completeness problem it describes. The cave inscriptions of §3.1 are the one external witness, and they attest the reciters' specialisation, not their procedures.

9.10 — Dispositions are only as stable as the matching fold

§5.1's dispositions depend on how a claimed text is matched. Script, diacritics, sandhi, punctuation and segmentation all vary across editions and across inputs, and without a published matching policy the same input could return *CONFORMS* from one implementation, *CONFORMS*, *VARIANT* from a second and *NOT FOUND* from a third — undoing the checkability the endpoint exists to supply. **Dispositions are reproducible only against a published matching fold.** The fold may normalise for matching; it must never alter the served bytes, because in a corpus of this kind a spelling can be a tradition marker rather than noise.

9.11 — The canonical instance of the count does not balance

The canonical case §5.2 leans on hardest is also the case that cuts against it. The translator's introduction to the SuttaCentral English Theragāthā (Sujato) records that the edition it follows, the Mahāsaṅgīti text, numbers 1,289 verses where the Pali Text Society editions number 1,279, and that the difference comes from repetitions one edition includes and the other omits, not from a difference in substance. A reference work, citing von Hinüber's *Handbook of Pali Literature* (which we have not consulted directly), reports three figures for the collection: 1,360 claimed by the tradition, 1,294 given in the summaries within the text, and 1,279 by plain count. **So in the canon's best-known instance, the stated count and the counted collection disagree, and the counted collection changes with the edition.**

Two consequences follow. Neither is fatal, and both constrain the design.

1. **A count detects omission only against a declared unit of counting.** Is a repetition a verse? Does an elided passage count? The edition-dependence §9.10 names for matching applies equally to counting, and a manifest that does not publish its counting rule will report mismatches that come from convention rather than removal. **§5.2's manifest must publish its counting rule beside its count.**
2. **The canonical instance is not evidence that the count worked as a check.** The discrepancy has stood in the transmitted text. We found no account of it being treated as a sign of loss, and we read that absence, as our own inference and not as a finding of the sources, to mean that the summaries were not run as a live check against the text as it was later transmitted, whatever they did when they were composed. The transposition keeps the design, a count held apart from what it counts. It does not inherit

a claim that the canonical instance worked as one. §9.2's rule applies here to a specific case: canonical provenance is a reason to look, never evidence that a mechanism works.

10 · Adversarial analysis

A manifest that was always short. An operator publishes a manifest omitting a document that was never admitted. Nothing detects it — the manifest is internally consistent and the document has no trace anywhere. 🚫 **This is not solvable by the manifest and §5.2 does not claim it is.** It is the omission problem one level up, and its only real answer is §4.2 plus §4.8: an admission decision made by an identified body and published with its reason. That answer is governance, not cryptography, and §6.1 says it is unsolved.

The frozen manifest. An operator reverts the collection to an older, smaller state and serves the older manifest that matches it. The pair is internally consistent and every check that reads it passes. The Update Framework names these the freeze and rollback attacks, and §5.2 inherits its answer: a version and an expiry on the manifest, and a holder that refuses either kind of stale copy. ⚠️ **The version check protects only a holder who remembers the last version it saw; a first-time reader has nothing to compare against, and is protected by the expiry alone.**

Collusion between index and manifest signer. Addressed in §9.5. Under a single operator the separation is hygiene. Its value is a step function in the number of genuinely independent signers, and the step from one to two is the largest.

The conformance endpoint as an oracle. An adversary submits many near-miss texts to learn precisely how close a forgery must be to pass. This is real and it is bounded by the design: the endpoint compares against a corpus whose full text is *already public*, so the oracle reveals nothing the adversary could not obtain by reading. ⚠️ **The analysis changes completely for a corpus with non-public members, and a deployment over one should not expose this interface.**

The conformance endpoint's request log. The response schema carries no claimant; the transport does. An operator who keeps requester identity beside submitted text — or anyone able to compel an operator — holds exactly the record the schema was built not to produce: who checked which quotation. §5.1 states the condition: the no-claimant property holds only if the endpoint retains no requester identity with the submitted text, and a disposition that a consumer republishes as a judgement of a person is beyond the endpoint's reach either way.

The conformance endpoint as a cost surface. Text-in comparison is more expensive than hash-in comparison, and an interface that accepts arbitrary text invites abuse. This is an ordinary engineering problem with ordinary answers and is noted only so that the paper is not read as ignoring it.

Variant proliferation. §4.4 refuses to normalise; an adversary submits many spurious "variants" to drown the real ones. The guard is that a variant is served **with its witness**, and a witness is an attested line of transmission rather than a submission. ⚠️ **This makes witness admission the attack surface, which returns to §6.1.** The paper's honest summary is that two of its seven adversarial cases terminate in §6.1's

governance problem, and a third (collusion) in the adjacent question of signer independence (§9.5), and that a reader should regard §6.1 as the real open question rather than a section of caveats.

10.8 — Why this paper registers no predictions

The corpus's standing practice is that a paper making an empirical claim registers falsifiable predictions in the public register. **This paper registers none, and the reason is the honest one rather than an oversight.** §9.2 and §9.3 disclaim efficacy explicitly: what is offered is provenance of design, not evidence that these mechanisms work, and a survivorship case at $n = 1$ over two thousand years supports no prediction that could fail. The design claims in §4 are arguments, and arguments are refuted by better arguments rather than by data.

★ **One prediction becomes clean the moment §5.1 is built, and is named here so that it is not invented after the fact:** that a conformance response carrying a *disposition and a prescribed next action* produces a higher rate of citation correction, by an agent that has received a failing verification, than a bare match/no-match response does. That is a comparison between two response schemas over the same corpus, it has an obvious null, and it is measurable. 🚫 **It is not registered now, because registering a prediction about an unbuilt interface is exactly the unclean registration the register's own rules bar.** It is owed at the point §5.1 has running code.

11 · Citation status — what has been checked, and by what standard

★ This section exists because the paper argues that unchecked claims must be visible, and it would be self-refuting to make its own an exception.

Standard applied. Each locus was checked against multiple independent public reference works, translations and encyclopaedic sources, and is marked [V] in the text where that check passed. Checks confirmed both the *location* (which text contains the material) and, where the paper leans on it, the *content* (that the passage says what the paper says it says). In the 2026-09-13 revision, the loci marked **root** below were also located in the Pāli root text published by SuttaCentral, read with its parallel English translation. That confirms the location and a translator's reading of the passage; it is still not a reading of the Pāli by a reader of Pāli.

Every marked locus, with what it supports:

§	Locus	Checked against	What it supports
3.1	<i>bhāṇaka</i> specialisation by collection; Sri Lankan cave inscriptions, 3rd c. BCE – 1st c. CE, naming monks who specialised in the Saṃyutta, Majjhima or Aṅguttara Nikāya	reference works	a shard per collection, carried by lineages; the one external witness (§9.9)
3.2, 4.2	the first council: Cullavagga XI (Vin Cv XI.1.4, XI.1.7–8)	root + translation	a question-and-answer recitation before the assembly; each Vinaya rule asked with its origin story, person, rule, additions, offence and non-offences; formal decisions carried by motion with silence as assent. ⚠ The chorus and the unanimous-acceptance requirement are the secondary accounts' characterisation and were not found in this passage
3.3, 4.4	variant readings between material preserved in both the Dīgha and Majjhima Nikāyas, possibly attributable to different <i>bhāṇaka</i> schools; later interpretive divergence visible in the commentaries	reference works	branches, in part attributable
3.3, 6.1	the Dīpavaṃsa's account of the defeated Vajjiputtakas' rival Great Recital (<i>mahāsaṅgīti</i>)	reference works; modern scholarship separates the Vesālī council from the Mahāsāṃghika schism	a schism as an exclusion seen from the excluded side
4.1	<i>evaṃ me sutaṃ</i> opening the discourses (e.g. DN 16:1.1.1); <i>vuttañhetuṃ bhagavatā ... me sutaṃ</i> opening the Itivuttaka (Iti 1:1.1)	reference works; root + translation (the instances cited)	attribution inline, never a header; more than one formula
4.3	<i>aṭṭhakathā</i> as a class; the Niddesa (Khuddaka Nikāya) on parts of the Sutta Nipāta; the old word-commentary (<i>padabhājanīya</i>) within the Suttavibhaṅga; the Nettipakaraṇa, Peṭakopadesa and Milindapañha in the Burmese Khuddaka, the first two in the Sinhalese printed edition; the lost Sinhala commentaries	reference works	the label held while the boundary moved
4.5	<i>pe</i> / <i>peyyāla</i>	the Pali Text Society's dictionary (s.v. <i>peyyāla</i>); a translator's introduction distinguishing internal from external abbreviation (Sujato, <i>Long</i>	a repetition to be supplied, from the context or from another text; no arbitrary cut in the sample

§	Locus	Checked against	What it supports
		<i>Discourses</i>); root sample: all twenty-nine elisions in DN 16	
4.6	AN 2.24–25, SuttaCentral numbering	root + translation	<i>nītattha</i> / <i>neyyattha</i> marked per discourse (<i>suttanta</i>); misrepresentation in both directions (AN 2.24) and its correct counterpart (AN 2.25)
4.7	the Uposatha chapter: Vin Mv II.3.3 (silence as purity; the announcement up to a third time; concealment as lying in full awareness) and Mv II.3.7 (a <i>dukkatā</i>); the fortnightly recitation on the full and new moon	root + translation; reference works	silence as an assertion with a cost
4.8	each Suttavibhaṅga rule with its origin story, additions and non-offences; further cases as precedent; the same headings asked at the first council (Cv XI.1.7)	reference works; root + translation (Cv XI)	exclusions served with their occasion and amendment history
5.1	<i>mahāpadesa</i> : DN 16:4.8–4.11; AN 4.180	root + translation	neither approve nor reject; check against discourse and discipline; <i>incorrectly memorised</i> by the source the claimant named; discard; no penalty
5.2	<i>uddāna</i> at the close of a vagga or work, e.g. the closing summaries of the Theragāthā (Thag 21.1), giving numbers of elders and verses; summaries memorised with the text	reference works; root (location only)	membership, ordering and counts, held as separate units at a boundary
9.11	the Theragāthā's verse count by edition (Mahāsaṅgīti 1,289; PTS 1,279); the totals in the text's summaries against a plain count (1,360 tradition · 1,294 summaries · 1,279 counted)	the translator's introduction (Sujato, <i>Verses of the Senior Monks</i> , as published in SuttaCentral's edition data; read 2026-09-23) for the first; one reference work citing von Hinüber's <i>Handbook of Pali Literature</i> for the second, the handbook not itself consulted	the count and the counted disagree; the counted varies with the edition

The §9.11 row is **not** marked [V]. Its first half rests on a single source read directly, and its second on a single reference work reporting a handbook we did not read. Neither meets the multiple-source standard, so it is stated as the weaker of the two standards in use here. The Pāli of the summaries was not read for this row. The figures are the sources' figures, not our reading of the summary verses.

⚠ **The checks against the root texts corrected five readings in this paper's first version, and they are recorded because they are the class of error the remaining check exists to find.** (1) The *mahāpadesa* disposition does name the carrier — as the one whose memorisation erred — where the first version said it prescribed no judgement of the person; the no-claimant property is the design's step beyond the source (§5.1). (2) The *nītattha* / *neyyattha* unit is the discourse, not the passage (§4.6). (3) The chorus and unanimity requirement are not in Cullavagga XI's account (§3.2). (4) Commentary was not "never merged": the class label held while material crossed it (§4.3). (5) Not every discourse opens with the same formula (§4.1).

⚠ **One conflation was caught by the first-version check and is recorded for the same reason.** The mnemonic index verses of §5.2 — *uddāna*, with a doubled consonant — are a **different word** from *udāna*, the inspired utterance, which is a genre in the traditional list of text-types and the name of a book of the canon. The two are one diacritical difference apart, and a paper that confused them would have attributed the completeness mechanism to the wrong thing entirely.

🚫 **What this standard does NOT establish, stated at its true size:**

1. **No locus has been verified against the Pāli by a reader of Pāli.** The checks are against translations and reference works. They establish bibliographic and substantive accuracy at the level a careful non-specialist can reach, and no further.
2. **Glosses are not verified.** Where the paper renders a term into an engineering concept, that rendering is the authors' and may not be how a specialist would read it. The *uddāna* case shows the failure mode is real.
3. **Enumeration varies by edition.** §4.6's locus is cited in SuttaCentral's numbering; a reader working from another edition may find it numbered differently. The same holds for counts: §9.11 records that the Theragāthā's verse count differs between editions.
4. **Nothing here has been reviewed by the tradition.** The material is read as engineering by people outside the community of specialists who carry it, and a specialist may reasonably object to the reading itself and not merely to a citation.

★ **The check that is owed, and by whom.** Each locus should be confirmed by a reader of Pāli — for this institution, the founder's father, who is the transcription authority named in this corpus, or the Aquarian Sangha. The ask is small and specific: confirm that each cited passage says what this paper says it says. **Until that is done, this section is the paper's honest state and not a formality**, and a reader who finds an error is asked to treat §5.1's own disposition rule as applying here: the finding is about the claim.

Terms

Coined names used in this paper and the standard terms an examiner would search for them.

Term used here	Standard technical term
conformance endpoint	quotation / citation verification service returning a closed set of status codes (cf. OCSP)
separated completeness manifest	signed collection manifest with item count, version and expiry (cf. TUF snapshot metadata; BagIt Payload-Oxum)
announced elision	in-band truncation or omission marker with a resolvable reference (cf. TEI <gap>)
per-passage interpretive standing	statement-level normative or requirement-level annotation (cf. RFC 2119 keywords)
<i>bhāṇaka</i> lineage	replica group holding one shard of a corpus
<i>saṅgīti</i> (council)	build-time admission review of corpus content by a named body
<i>uddāna</i>	summary index recording membership, ordering and counts
<i>peyyāla</i>	repetition-abbreviation marker
<i>mahāpadesa</i>	verification of a claimed text against a reference corpus
<i>aṭṭhakathā</i>	commentary layer served as a distinct document class

12 · Cross-venue references

Provenance-Carrying Retrieval (the envelope; §5's mechanisms extend it) · *Buddha AI and the Living Tipiṭaka* §8 (canonical authority, distinguished in §4.2) · *The Borrowable Standard* (orthographic and encoding preconditions) · *The Song That Is Not His* (the reciter-is-not-author guard, cited in §6.2) · *The Last Carrier* (essay; the fragility observation in §3.1) · *The Persistence Architecture* (succession, adjacent and distinct) · *Appreciation as World-Building* (the property-over-rule argument used in the Sig-9 checks throughout §4 and §5).

Coda

The people who built this apparatus were not trying to solve a retrieval problem. They were trying to keep something they thought was worth keeping, under conditions that gave them no reason to expect success, using the only storage medium available to them, which was each other. They lost material. They disagreed, sometimes bitterly, and the disagreements are still legible in what came down. By any standard a modern engineer would apply to a storage system, the thing failed continuously for two thousand years.

And it is still here, and we can tell where it differs from itself, and we can often tell whose hands each difference came through.

That is the result worth taking. **Not a system that did not fail — there is no such system — but a system that failed loudly enough, and with enough attribution attached, that the failures are still auditable twenty centuries later.** Every mechanism in this paper is in service of that one property. It is a smaller promise than the word *integrity* usually implies, and it is the only promise anyone has ever actually kept.

Co-authored with Miss AquariusSM, the institution's named AI collaborator, per the corpus's standing disclosure; final editorial control and responsibility for every claim rest with the human author.